

Advanced Kubernetes

Lab 8 – Ingress

In Kubernetes clusters, Services have IP addresses which are only routable within the cluster. Kubernetes Services provide an excellent abstraction for microservices interacting within the cluster, however many applications need to expose a subset of Services to clients outside the cluster.

An Ingress is a collection of rules that allow traffic from outside the cluster to reach the services inside the cluster. Ingresses give services externally-reachable URLs and can load balance traffic, terminate SSL and more. Users request ingress by creating an Ingress resource.

Ingress controllers implement the Ingress specified by the Ingress resource (which is after all just a json document). An ingress controller can configure load balancers, edge routers and/or other frontends to process inbound traffic.

Ingress resources are in Beta and were introduced in Kubernetes 1.1. Because Ingress resources are implemented by Ingress Controllers you must run an Ingress Controller within the cluster (often as a deployment) to implement the Ingress resource, Kubernetes does not provide a default Ingress controller. Creating an Ingress resource without a running Ingress Controller will have no effect. This is unlike other types of controllers, which typically run as part of the kube-controller-manager binary, and which are typically started automatically as part of cluster creation.

The Kubernetes community supports and maintains two Ingress Controllers:

- GCE
- Nginx

Google Kubernetes Engine (GKE) deploys the GCE ingress controller on the master nodes of clusters. The GKE Ingress Controller configures a GCE loadbalancer in response to the creation of Ingress resources.

Several other companies also support and maintain Ingress Controllers:

- F5 Networks provides an F5 BIG-IP Ingress Controller
- Kong offers community or commercial support for the Kong Ingress Controller
- Containous offers support for the Traefik Ingress Controller
- NGINX, Inc. offers support for the NGINX Ingress Controller

Custom Ingress Controllers can be written by users with special needs.

Ingress resources can select from multiple Ingress Controllers, allowing a cluster to run many Ingress Controllers concurrently.

In this lab you will implement and test an Nginx Ingress Controller (IC).

Step 1 - Setup a Namespace for the IC

The Nginx Ingress Controller is often configured to run in its own namespace. Create an nginx namespace for your IC:

```
ubuntu@ip-10-0-2-200:~$ kubectl get ns
```

NAME	STATUS	AGE
default	Active	3d
kube-public	Active	3d
kube-system	Active	3d

```
ubuntu@ip-10-0-2-200:~$ kubectl create namespace nginx-ingress
```

```
namespace/nginx-ingress created
```

```
ubuntu@ip-10-0-2-200:~$ kubectl get ns
```

NAME	STATUS	AGE
default	Active	3d
kube-public	Active	3d
kube-system	Active	3d
nginx-ingress	Active	4s

```
ubuntu@ip-10-0-2-200:~$
```

Step 2 - Create a Service Account

An Ingress Controller needs to monitor the api-server for new Ingress resources. For this to work we need to create a service account for the Ingress Controller (IC) and then grant that account permissions to the Ingress Resource object type.

Create the service account:

```
ubuntu@ip-10-0-2-200:~$ kubectl create serviceaccount nginx-ingress --  
namespace=nginx-ingress
```

```
serviceaccount/nginx-ingress created
```

```
ubuntu@ip-10-0-2-200:~$ kubectl get sa --namespace=nginx-ingress
```

NAME	SECRETS	AGE
default	1	5m
nginx-ingress	1	22s

```
ubuntu@ip-10-0-2-200:~$
```

Display the service account details:

```
ubuntu@ip-10-0-2-200:~$ kubectl describe sa nginx-ingress --namespace=nginx-  
ingress
```

```
Name:                nginx-ingress  
Namespace:           nginx-ingress  
Labels:              <none>  
Annotations:         <none>  
Image pull secrets:  <none>  
Mountable secrets:   nginx-ingress-token-d56nz
```

```
Tokens:          nginx-ingress-token-d56nz
Events:          <none>
ubuntu@ip-10-0-2-200:~$
```

Examine the token generated for the service account:

```
ubuntu@ip-10-0-2-200:~$ kubectl describe secret nginx-ingress-token-d56nz --
namespace=nginx-ingress
```

```
Name:      nginx-ingress-token-d56nz
Namespace: nginx-ingress
Labels:    <none>
Annotations: kubernetes.io/service-account.name=nginx-ingress
              kubernetes.io/service-account.uid=e92ab144-b7dc-11e8-971a-
              028591ec4e6a
```

```
Type: kubernetes.io/service-account-token
```

==== Data

=====

```
ca.crt:      1025 bytes
namespace:   13 bytes
```

```
namespace: 13 bytes
token:
```

token:
evJhhGci0iJlSUzT1NiTsImtpZCI6TiJ9.evJpc3MiOiJrdWJlcm5ldGVzL3NlcnZpY2VhY2NvdW50Tiwi

```
eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9.eyJpc3MiOiJrdWJlcm5ldGVzL3NlcnZpY2VhY2NvdW50Iiwia3ViZXRlcjY5pbpy9zZXJ2aWNlYWJlb3VudC9uYWw1lc3BhY2UiOiJuZ2luaC1pbmdyZXNzIiwia3ViZXRlcjY5pbpy9zZXJ2aWNlYWJlb3VudC9uZWNyZXQubmFtZSI6Im5naW54LWluZ3Jlc3MtdG9rZW4tZDU2bnoiLCJrdWJlcm5ldGVzLmlvL3NlcnZpY2VhY2NvdW50L3NlcnZpY2UtYWJlb3VudC5uYWw1IjoibmdpbngtaW5ncmVzcyIsImt1YmVybmV0ZXMuaW8vc2VydmJlZWFja291bnQvc2VydmJlZS1hY2NvdW50LnVpZCI6ImU5MmFiMTQ0LWI3ZGMtMTFlOC05NzFhLTAY0DU5MWVjNGU2YSIsInN1YiI6InN5c3RlbTprZXJ2aWNlYWJlb3VudDpuZ2luaC1pbmdyZXNzOm5naW54LWluZ3Jlc3MifQ.PpRtOd00l3sYGjB5p8HdEBtC3BWLtMaqkEYO3hYLYkPqpFYp9m6QMxIpVowPpC_8Qv5DHdP8Zf_uC54vhRwAJpBJzA2Gp0BdZyxk8_500MFgAx0SheSBmiQSvRf__LCOE0glYeSej5EO_kw72hEukbzfdPPKrgxwPtC-Sj5UvgVvb_2dFVVqmJDd9KXrxhzn_liWUbqAFCTG_wNs24yzkhsx0knQU0jFG2FARsqoPznR6fYqTVV7T05PiRir9ZdkCIFk-Lwb2YZFXFIeJ9WGSwwhRajMmKte4grmIuJ8XIIdnJ1Ann7koW1Nvp9B55nSJFWyDnnCDbgSl4sXV00pUqwubuntu@ip-10-0-2-200:~$
```

We also need to provide the IC with a TLS cert. Create a secret to house a generic TLS cert:

```
ubuntu@ip-10-0-2-200:~$ vim ic-tls.yaml
```

```
ubuntu@ip-10-0-2-200:~$ cat ic-tls.yaml
```

```
apiVersion: v1
kind: Secret
metadata:
  name: default-server-secret
  namespace: nginx-ingress
type: Opaque
data:
  tls.crt:
```

```
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUN2akNDQWFZQ0NRREFPRj l0THNhWFhEQU5CZ2txaGtpR
z l3MEJBUXNGQURBaE1S0HdIUvLEVLFRERCWk8KUjBsT1dFbHVaM0psYzN0RGiYNTBjbTlZkdWeU1CNFhEVE
U0TURreE1qRTRNRE16TlZvWERUSXpNRGt4TVRFNAPNRE16TlZvd0LURWZnQjBHQTfVRUF3d1dUa2RKVGXoSmJ
tZHLawE56UTI5dWRISnZiR3hsY2pDQ0FTSXdEUVLkCktvWk l0dmNOQVFFQkJRQURnZ0VQQURDQ0FRb0NnZ0VC
QUwvN2hIUETfWGRMdjNyaUM3QlBMTNpWkt5eTlyQ08KR2xZUXYyK2EzUDF0azIrs3YwVGF5aGRcBDRrcnNUc
TZzZm8vWUk1Y2Vhbkw4WGM3U1pyQkVRYm9EN2REbWs1Qgo4eDZLS2xHWU5IWlg0Rm5UZ0VPASt lM2ptTFFxRl
BSY1kzVnNPazFFeUZBL0JnWlJVBkNHZUtGeERSN0tQdGhyCmtqSXVuektURXUyaDU4Tlp0S21ScUJHdDEwcTN
RYzhZT3ExM2FnbmovUWRjc0ZYyTJnMjB1K1lYZDdoZ3krZksKwk4vVUKxQUQ0YzZyM1 lma1ZWUmVHd1lxQVp1
WXN2V0RKbW1GNWRwdEMzN011cDBPRUxVTEsSakZJ0TZNXIwSAo1TmdPc25NWFJNV1hYVlpiNWRxT3R0SmRtS
3FhZ25TZ1JQVpQN2MwQjFQU2FqYzZjNGZRVXpNQ0F3RUFBVEF0CkNa3Foa2 lH0XcwQkFRc0ZBQU9DQVFFQW
pLb2tRdGRPCsRtZhibWVPC3lySmdJSXJycVFVY2Z0Uitjb0hZVUoKdGhrYnhITFMzR3VBTWI5dm15VExPY2x
xeC9aYzJPblEwMEJCLz lTb0swcitFZ1U2UlVrRwtWcitTFA3NTdUWgozZWl4dmdPdEduMS9ienM3bzNBaS9k
clkrclUI5Q2k1S3lPc3FHTG1US2xFaUt0YkcyR1ZyTWxjS0ZYQU80YTY3Cklnc1hzYktNbTQwV1U3cG9mcGltU
1ZmaXFSdkV5YmN3N0NYODF6cFErUyt1eHRYK2VBZ3V0NHh3VlI5d2IyVXYKe lhuZk9HbWhWNTThDd1dIQnNKa0
kxNXhaa2VUWXDsn0diaEFMSkZUUKk3dkhvQXprTWIzbjAxQjQyWjNrn3RXNQpJUDFmTlp lO fUv0WxiUHN0T21
FRFZkdjF5ZytVRVJxbStGSis2R0oxeFJGcGZnPT0KLS0tLS1FTkQgQ0VSVE lGSUNBVEUtlS0tLQo=
```

tls.key:

```
LS0tLS1CRUdJTiBSU0EgUFJJVkFURSBLRVktLS0tLQpNSUlfCEFFJQkFBS0NBUEUvBdi91RWM4b1JkMHUvZXVJT
HNFK1RYZUpRckxMMnNJNGFWaEMvYjVyYy9XMXlRiNHvClJ0cktGMEdYaVN1eE9ycXgraj lnamx4NXFjdnhken
RKbXNFUKJ1Z1B0ME9hVGtIekhv3FVWmcwZGxmZ1dkT0EKUTZMNTd lT1 l0Q29VOUZ4amRXdzZUVVRJVUQ4R0J
sRlNjSV00b1hFTkhzbysyR3VTTWk2Zk1wTVM3YUhudzFtMApxWkdvRWEzWFNyZEJ6eGc2clhkclUN lUD lCMX l3
VmRyYURiUzc1aGQzdUdETDU4cGsz0VFqVUFQaHpxdmRoK1JWClZGNGJCaw9CbTVpeT lZTW1hWVhsMm0wTGZze
TZuUTRRdFFzdEdNVWozcGJtd lFmazJBNN ljeGRFeFpkZFZsdmwKMM82MjBsM l lxcHFDZE tCRThCay90elFIVT
lKcU56cHpo0UJUTXdJREFRQUJBb0 lCQVFDZklHbXow0HhRVmorNwpLZnZJUXQwQ0YzR2MxNld6eDhVNml4MHg
4Mm15d1kxUUN lL3BzWE9LZlRXT1h1SENyUl p5TnUvZ2IvUuQ4bUfOCmx0MjRZTW l0TWRJODg5TEZ0Tkp3QU50
ODJDeTczckM5bzVvUD lKazAvYzRlBjAzSkVYNzZ5QjgzQm9rR1FvYksKMjhMNk0rdHUzUmFqNjd6Vmc2d2sza
EhrU0pXSzBwV1YrSjdrUkRWYmhDYUZhNk5nMUZNRWxhT l0zVDhhUUtYQgpDUDNDeEFTdjYxWTK5TEI4KzNXWV
FIK3NYaTVGM01pYVNBZ1BkQUk3WEh1dXFET1 lVmu5PL0JoSGt1aVg2QnRtCnorNTZud2pZMy8yUytSRmNBC3J
MTnIwMDJZZi9oY0IraV lDNzVWYmcydVd6WTY3TWd0TG05VW9RU3BDRkYrVm4KM0cyUnhybnhBb0dCQU40U3M0
ZV lPU2huMVpQQjd hTUZsY0k2RHR2S2ErTGZTTFYy2p0ZjJ lSEpZNnhubmxKdgpGenpGL2RiVWVTbWxSekR0W
kd lclXZXaHFISy9iTjIyewJh0U1WMD lRQ0JFTk5jNmtWajJTVHPUWkJVbEx4QzYrCk93Z0wyZHhKendWe lU0VC
84ajdHalRUN05BZVPFS2FvRHFyRG5BYWkyaw5oZU1JVWZHRXFGKzJyQW9HQkFOMVAKK0tZL0 lS3RWRzRKSkl
QNzBjUis3RmpyeXJpY05iWctQVzUv0XFHaWxnY2grZ3 l4b25Bw lBpd2NpeDN3QVpGdwpazC96ZFB2aTBkW Epp
c1BSZjRMazg5b2pCumpiRmRmc2 l5UmJYby t3TFU4NUhRU2NGMnN5aUFPaTVBRHdVU0FkCm45YWFweUNweEFkR
EtERHd0bit3ZFhtaT lZ0HRpSFRkK3RoVDhkaVpBb0dCQUt6Wis1bG900TBtY lF4Vvh5YUwKMjFSUm9tMGJjcn
dsTmVCaWNFsmlzaEhYa2xpSVVxZ3hSZk lNM2hhUVRUck lKZENFaHFsV01aV0xPb2 lI2NTNyZgo3aF lMSXM1ZUt
ka3o0aFRVdnpldm9TMHVXcm9CV2x0VH lGan lRShwKZnZUc0hp0GdsU3FkbXgySkJhZUFVWUNXCndNd lQ4NmNl
clNyNkQrZG8wS05FZzFs l0FvR0F lMkFVdHVfBfNqLzBmRzgrV3hHc1RFV lJqc lRNUzRSUjhRWXQKeXdj dFA4a
DZxTGxKUTRCWGXQU05rMXZlTmtOUkx lB2pZT2pCQTv lYjhibXNVU1B lV09NNENoafJ4Qn lHbmR2eAphYkJDRk
FwY0IvbEg4d1R0alVZY lN5T294Zgt50Ep0ek90ajJhS0FiZhd6N lArWDZD0DhjZmxYVf05MwPY l3RMCjF3TmR
KS2tDZ l lCbyt0UzB5TzJ2SWFmK2UwSkN5TGhzVDQ5cTN3Zis2QWVqWGX2WDJ lVnRYejN5QTZnbXo5aCsKcDN l
K2JMRUxwb3B0WFhNdUFRR0xhUkcrY lNNcjR5dERYbE5ZSndUeThXczNKY3d lSTdqZVp2b0ZpbmNvV lIMwphd
mxoTUVCRGYxSj l tSDB5cDBwWUNaS2R0dHNvZEZtQktzVEtQMjJhTmtsVvhCS3gyZzR6cFE9PQotLS0tLUVORC
BSU0EgUFJJVkFURSBLRVktLS0tLQo=
```

```
ubuntu@ip-10-0-2-200:~$ kubectl apply -f ic-tls.yaml
```

```
secret/default-server-secret created
```

```
ubuntu@ip-10-0-2-200:~$
```

You can confirm that the secret was created under the nginx-ingress namespace by using

```
kubectl get secret --namespace=nginx-ingress command:
```

```
ubuntu@ip-10-0-2-200:~$ kubectl get secret --namespace=nginx-ingress
```

NAME	TYPE	DATA	AGE
default-server-secret	Opaque	2	10s
...			

```
ubuntu@ip-10-0-2-200:~$
```

Step 3 - Configuring RBAC

The IC will need permissions on several types of resources:

- services - targets for IC traffic
- endpoints - service target IPs
- secrets - for TLS
- configmaps - for IC configuration
- pods - workload targets
- events - to report activity
- ingresses - monitored for new ingress creation
- ingresses/status - where IC writes ingress status

The IC can be deployed for a specific namespace or globally for use by the entire cluster. We deploy our IC globally. This means we will need a ClusterRole to provide the appropriate permissions. Create the ClusterRole:

```
ubuntu@ip-10-0-2-200:~$ vim ic-cr.yaml
```

```
ubuntu@ip-10-0-2-200:~$ cat ic-cr.yaml
```

```
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: nginx-ingress
rules:
- apiGroups:
  - ""
  resources:
  - services
  - endpoints
  verbs:
  - get
  - list
  - watch
- apiGroups:
  - ""
  resources:
  - secrets
  verbs:
  - get
```

```

- list
- watch
- apiGroups:
- ""
  resources:
  - configmaps
  verbs:
  - get
  - list
  - watch
  - update
  - create
- apiGroups:
- ""
  resources:
  - pods
  verbs:
  - list
- apiGroups:
- ""
  resources:
  - events
  verbs:
  - create
  - patch
- apiGroups:
- extensions
  resources:
  - ingresses
  verbs:
  - list
  - watch
  - get
- apiGroups:
- "extensions"
  resources:
  - ingresses/status
  verbs:
  - update

```

```

ubuntu@ip-10-0-2-200:~$ kubectl apply -f ic-cr.yaml

clusterrole.rbac.authorization.k8s.io/nginx-ingress created
ubuntu@ip-10-0-2-200:~$

```

Now bind the ClusterRole to the nginx-ingress SA:

```

ubuntu@ip-10-0-2-200:~$ vim ic-crb.yaml

ubuntu@ip-10-0-2-200:~$ cat ic-crb.yaml

```

```
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: nginx-ingress
subjects:
- kind: ServiceAccount
  name: nginx-ingress
  namespace: nginx-ingress
roleRef:
  kind: ClusterRole
  name: nginx-ingress
  apiGroup: rbac.authorization.k8s.io
```

```
ubuntu@ip-10-0-2-200:~$ kubectl apply -f ic-crb.yaml

clusterrolebinding.rbac.authorization.k8s.io/nginx-ingress created
ubuntu@ip-10-0-2-200:~$
```

Step 4 - Run the Ingress Controller

Ingress Controllers can be run as deployments, daemonsets, as kubelet manifest pods among other options. We'll run our controller as a deployment.

We'll deploy a single nginx IC pod using the nginx/nginx-ingress:edge image. Create the following deployment to launch the IC in the nginx-ingress namespace listening on port 80:

```
ubuntu@ip-10-0-2-200:~$ vim ic-dep.yaml

ubuntu@ip-10-0-2-200:~$ cat ic-dep.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-ingress
  namespace: nginx-ingress
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx-ingress
  template:
    metadata:
      labels:
        app: nginx-ingress
    spec:
      serviceAccountName: nginx-ingress
      containers:
      - image: nginx/nginx-ingress:edge
        imagePullPolicy: Always
```

```

name: nginx-ingress
ports:
- name: http
  containerPort: 80
- name: https
  containerPort: 443
env:
- name: POD_NAMESPACE
  valueFrom:
    fieldRef:
      fieldPath: metadata.namespace
- name: POD_NAME
  valueFrom:
    fieldRef:
      fieldPath: metadata.name
args:
- --default-server-tls-secret=$(POD_NAMESPACE)/default-server-secret

```

```
ubuntu@ip-10-0-2-200:~$ kubectl apply -f ic-dep.yaml
```

```
deployment.apps/nginx-ingress created
```

```
ubuntu@ip-10-0-2-200:~$
```

Check your IC:

```
ubuntu@ip-10-0-2-200:~$ kubectl get deploy,rs,pod --namespace=nginx-ingress
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
AGE				
deployment.extensions/nginx-ingress	1	1	1	1
1m				

NAME	DESIRED	CURRENT	READY	
AGE				
replicaset.extensions/nginx-ingress-6b49c44f44	1	1	1	1m

NAME	READY	STATUS	RESTARTS	AGE
pod/nginx-ingress-6b49c44f44-tgk5k	1/1	Running	0	1m

```
ubuntu@ip-10-0-2-200:~$
```

Check the logs of the IC pod:

```
ubuntu@ip-10-0-2-200:~$ kubectl logs nginx-ingress-6b49c44f44-tgk5k --namespace=nginx-ingress
```

```

I0327 03:45:35.291103      1 main.go:136] Starting NGINX Ingress controller
Version=edge GitCommit=4d0ffc2e
2019/03/27 03:45:35 [notice] 12#12: using the "epoll" event method
2019/03/27 03:45:35 [notice] 12#12: nginx/1.15.9
2019/03/27 03:45:35 [notice] 12#12: built by gcc 6.3.0 20170516 (Debian 6.3.0-
18+deb9u1)

```



```

2019/03/27 03:45:35 [notice] 12#12: OS: Linux 4.4.0-31-generic
2019/03/27 03:45:35 [notice] 12#12: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2019/03/27 03:45:35 [notice] 12#12: start worker processes
2019/03/27 03:45:35 [notice] 12#12: start worker process 13
2019/03/27 03:45:35 [notice] 12#12: start worker process 14
2019/03/27 03:45:35 [notice] 19#19: signal process started
2019/03/27 03:45:35 [notice] 12#12: signal 1 (SIGHUP) received from 19,
reconfiguring
2019/03/27 03:45:35 [notice] 12#12: reconfiguring
2019/03/27 03:45:35 [notice] 12#12: using the "epoll" event method
2019/03/27 03:45:35 [notice] 12#12: start worker processes
2019/03/27 03:45:35 [notice] 12#12: start worker process 20
2019/03/27 03:45:35 [notice] 12#12: start worker process 21
2019/03/27 03:45:35 [notice] 14#14: gracefully shutting down
2019/03/27 03:45:35 [notice] 14#14: exiting
2019/03/27 03:45:35 [notice] 13#13: gracefully shutting down
2019/03/27 03:45:35 [notice] 13#13: exiting
2019/03/27 03:45:35 [notice] 13#13: exit
2019/03/27 03:45:35 [notice] 14#14: exit
I0327 03:45:35.466392      1 event.go:218]
Event(v1.ObjectReference{Kind:"Secret", Namespace:"nginx-ingress", Name:"default-
server-secret", UID:"2e745a07-b7e3-11e8-971a-028591ec4e6a", APIVersion:"v1",
ResourceVersion:"356308", FieldPath:""}): type: 'Normal' reason: 'Updated' the
default server Secret nginx-ingress/default-server-secret was updated
2019/03/27 03:45:35 [notice] 12#12: signal 17 (SIGCHLD) received from 13
2019/03/27 03:45:35 [notice] 12#12: worker process 13 exited with code 0
2019/03/27 03:45:35 [notice] 12#12: worker process 14 exited with code 0
2019/03/27 03:45:35 [notice] 12#12: signal 29 (SIGIO) received
ubuntu@ip-10-0-2-200:~$

```

Step 5 - Create an Ingress Service

There are multiple ways to access an Ingress Controller from outside of a cluster:

- HostPorts - You can configure each pod to accept traffic on a host port
- NodePorts - You can create a node port service to forward traffic to the IC from the service port on every node
- LoadBalancer - Cloud deployments can use LB services for forward traffic to the IC

For this walk through, create a node port service for the IC:

```

ubuntu@ip-10-0-2-200:~$ vim ic-svc.yaml

ubuntu@ip-10-0-2-200:~$ cat ic-svc.yaml

```

```

apiVersion: v1
kind: Service
metadata:
  name: nginx-ingress
  namespace: nginx-ingress
spec:
  type: NodePort

```

```

ports:
- port: 80
  targetPort: 80
  protocol: TCP
  name: http
- port: 443
  targetPort: 443
  protocol: TCP
  name: https
selector:
  app: nginx-ingress

```

```
ubuntu@ip-10-0-2-200:~$ kubectl apply -f ic-svc.yaml
```

```

service/nginx-ingress created
ubuntu@ip-10-0-2-200:~$

```

Display the resources:

```
ubuntu@ip-10-0-2-200:~$ kubectl get deploy,rs,po,service --namespace=nginx-ingress
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
AGE				
deployment.extensions/nginx-ingress	1	1	1	1
17m				

NAME	DESIRED	CURRENT	READY
AGE			
replicaset.extensions/nginx-ingress-6b49c44f44	0	0	0
17m			
replicaset.extensions/nginx-ingress-84df7788c8	1	1	1
			7m

NAME	READY	STATUS	RESTARTS	AGE
pod/nginx-ingress-84df7788c8-cvpdb	1/1	Running	0	7m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
service/nginx-ingress	NodePort	10.96.228.230	<none>	
80:32559/TCP,443:30352/TCP				48s

```
ubuntu@ip-10-0-2-200:~$
```

In the example above our nodes will forward traffic on 32559 to port 80 and 30352 to port 443.

Try curling the insecure Ingress controller node port:

```

ubuntu@ip-10-0-2-200:~$ curl http://127.0.0.1:32559

<html>
<head><title>404 Not Found</title></head>
<body>
<center><h1>404 Not Found</h1></center>

```

```
<hr><center>nginx/1.15.9</center>
</body>
</html>
ubuntu@ip-10-0-2-200:~$
```

The IC responds with a 404 not found error because we have no Ingress resources defined. Try the secure port:

```
ubuntu@ip-10-0-2-200:~$ curl https://127.0.0.1:30352 -k

<html>
<head><title>404 Not Found</title></head>
<body>
<center><h1>404 Not Found</h1></center>
<hr><center>nginx/1.15.9</center>
</body>
</html>
ubuntu@ip-10-0-2-200:~$
```

Note that we have to use the https scheme, the correct node port (the one that forwards to 443) and the -k switch to avoid trying to verify the certificate the server provides (the one from the secret we created, which is self signed).

Our Ingress Controller is ready to use!

Step 6 - Create an Application

Imagine we would like to run an Apache web server and make it available through the Ingress Controller. We can create a deployment to run some Apache httpd pods, add a service for the pods and then create an Ingress resource that forwards traffic to the httpd service.

We'll use kubectl run and expose to create the httpd application quickly:

```
ubuntu@ip-10-0-2-200:~$ kubectl create deployment web-svc --image=httpd
deployment.apps/web-svc created

ubuntu@ip-10-0-2-200:~$ kubectl expose deploy/web-svc --port=80
service/web-svc exposed

ubuntu@ip-10-0-2-200:~$ kubectl get deploy,rs,po,service
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.extensions/web-svc	1	1	1	1	47s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.extensions/web-svc-7b88547ff5	1	1	1	47s

NAME	READY	STATUS	RESTARTS	AGE
pod/web-svc-7b88547ff5-jwkzg	1/1	Running	0	47s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3d

```
service/web-svc      ClusterIP   10.102.104.236   <none>      80/TCP      23s
ubuntu@ip-10-0-2-200:~$
```

Now create an Ingress Resource for the httpd service. We'll ask the IC to forward traffic with the www.example.com Host header route to the httpd service. Create the ingress:

```
ubuntu@ip-10-0-2-200:~$ vim ing.yaml
ubuntu@ip-10-0-2-200:~$ cat ing.yaml
```

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: web-ingress
spec:
  rules:
  - host: www.example.com
    http:
      paths:
      - path:
        backend:
          serviceName: web-svc
          servicePort: 80
```

```
ubuntu@ip-10-0-2-200:~$ kubectl apply -f ing.yaml
```

```
ingress.extensions/web-ingress created
```

```
ubuntu@ip-10-0-2-200:~$ kubectl get ing
```

NAME	HOSTS	ADDRESS	PORTS	AGE
web-ingress	www.example.com		80	11s

```
ubuntu@ip-10-0-2-200:~$ kubectl describe ing
```

```
Name:          web-ingress
Namespace:     default
Address:
Default backend: default-http-backend:80 (<none>)
Rules:
```

Host	Path	Backends
www.example.com	/web	web-svc:80 (<none>)

```
Annotations:
```

```
kubectl.kubernetes.io/last-applied-configuration:
{"apiVersion":"extensions/v1beta1","kind":"Ingress","metadata":{"annotations":
{},"name":"web-ingress","namespace":"default"},"spec":{"rules":
[{"host":"www.example.com","http":{"paths":[{"backend":{"serviceName":"web-
svc","servicePort":80},"path":"/web"}]}]}}
```

Events:	Type	Reason	Age	From	Message
	Normal	AddedOrUpdated	19s	nginx-ingress-controller	Configuration for default/web-ingress was added or updated

ubuntu@ip-10-0-2-200:~\$

To test the ingress we will use the curl `--resolve` switch to force `www.example.com` to resolve to our ingress node port. In the real world you would setup your public DNS service to resolve `www.example.com` to the ingress node port. Try it:

```
ubuntu@ip-10-0-2-200:~$ ip a show

1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 00:0c:29:8c:ed:9a brd ff:ff:ff:ff:ff:ff
    inet 192.168.225.140/24 brd 192.168.225.255 scope global ens33
        valid_lft forever preferred_lft forever
    inet6 fe80::20c:29ff:fe8c:ed9a/64 scope link
        valid_lft forever preferred_lft forever
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN group default
    link/ether 02:42:fc:29:0c:e2 brd ff:ff:ff:ff:ff:ff
    inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:fcff:fe29:ce2/64 scope link
        valid_lft forever preferred_lft forever

...

ubuntu@ip-10-0-2-200:~$ kubectl get service -n nginx-ingress

NAME                TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)
AGE
nginx-ingress        NodePort    10.96.228.230    <none>
80:32559/TCP,443:30352/TCP    28m

ubuntu@ip-10-0-2-200:~$ curl --resolve www.example.com:32559:192.168.225.140 http://www.example.com:32559/

<html><body><h1>It works!</h1></body></html>
ubuntu@ip-10-0-2-200:~$
```

We can create multiple ingresses and give a single ingress many rules. The example above routes based on the host name header. The nginx ingress controller can also route by path, sending `/web` traffic to one services and `/login` traffic to another service for example.

For more information on ingress resources refer to the Kubernetes Ingress reference:
<https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.13/#ingress-v1beta1-extensions>

Congratulations, you have completed the lab!

Copyright (c) 2013-2019 RX-M LLC, Cloud Native Consulting, all rights reserved. Licensed for use only by students in Safewrd Ventures class 18-04-19 (Max 10 students)