

# Kubernetes

## Lab 2 – Kubernetes Exploration

Kubernetes clusters track and manage objects of various “kinds”. Applications make use of four kinds of objects in particular:

- Pods – groups of containers deployed as a unit
- Replica Sets – sets of pods defined by a template which the Controller Manager replicates across the cluster
- Deployments – a rollout strategy for pods and replica sets
- Services – end points used to distribute requests to one of a pod’s replicas

Thus basic Kubernetes applications consist of pods, which implement the application functionality; replica sets, which ensure pods are always available; and Services which expose a dynamic set of pods to clients as a single endpoint. deployments describe how to launch or upgrade a given application.

### 1. kubectl

The `kubectl` command provides a range of features we can use with Kubernetes. Run `kubectl` without arguments to get a list of the available commands.

```
ubuntu@ip-10-0-2-200:~$ kubectl
```

kubectl controls the Kubernetes cluster manager.

Find more information at: <https://kubernetes.io/docs/reference/kubectl/overview/>

#### Basic Commands (Beginner):

create	Create a resource from a file or from stdin.
expose	Take a replication controller, service, deployment or pod and expose it as a new Kubernetes Service
run	Run a particular image on the cluster
set	Set specific features on objects

#### Basic Commands (Intermediate):

explain	Documentation of resources
get	Display one or many resources
edit	Edit a resource on the server
delete	Delete resources by filenames, stdin, resources and names, or by resources and label selector

#### Deploy Commands:

rollout	Manage the rollout of a resource
scale	Set a new size for a Deployment, ReplicaSet, Replication Controller, or Job
autoscale	Auto-scale a Deployment, ReplicaSet, or ReplicationController

#### Cluster Management Commands:

certificate	Modify certificate resources.
cluster-info	Display cluster info
top	Display Resource (CPU/Memory/Storage) usage.
cordon	Mark node as unschedulable
uncordon	Mark node as schedulable
drain	Drain node in preparation for maintenance
taint	Update the taints on one or more nodes

#### Troubleshooting and Debugging Commands:

describe	Show details of a specific resource or group of resources
logs	Print the logs for a container in a pod
attach	Attach to a running container
exec	Execute a command in a container
port-forward	Forward one or more local ports to a pod
proxy	Run a proxy to the Kubernetes API server
cp	Copy files and directories to and from containers.
auth	Inspect authorization

#### Advanced Commands:

diff	Diff live version against would-be applied version
apply	Apply a configuration to a resource by filename or stdin
patch	Update field(s) of a resource using strategic merge patch
replace	Replace a resource by filename or stdin
wait	Experimental: Wait for a specific condition on one or many resources.
convert	Convert config files between different API versions

#### Settings Commands:

label	Update the labels on a resource
annotate	Update the annotations on a resource
completion	Output shell completion code for the specified shell (bash or zsh)

#### Other Commands:

api-resources	Print the supported API resources on the server
api-versions	Print the supported API versions on the server, in the form of "group/version"
config	Modify kubeconfig files
plugin	Provides utilities for interacting with plugins.
version	Print the client and server version information

#### Usage:

kubectl [flags] [options]

Use "kubectl <command> --help" for more information about a given command.

Use "kubectl options" for a list of global command-line options (applies to all commands).

ubuntu@ip-10-0-2-200:~\$

Take a moment to review available options. One useful subcommand is the global options, take a moment to review the output of `kubectl options`.

To use the `kubectl` command to control a remote cluster we must specify the cluster endpoint to `kubectl`. The `kubectl` command can be used to control several clusters from a single workstation. Clusters are given a name and

settings, including the IP address and port of the cluster API service.

To get configuration help issue the `kubectl help` subcommand.

```
ubuntu@ip-10-0-2-200:~$ kubectl help config
```

Modify kubeconfig files using subcommands like "kubectl config set current-context my-context"

The loading order follows these rules:

1. If the `--kubeconfig` flag is set, then only that file is loaded. The flag may only be set once and no merging takes place.
2. If `$KUBECONFIG` environment variable is set, then it is used as a list of paths (normal path delimitting rules for your system). These paths are merged. When a value is modified, it is modified in the file that defines the stanza. When a value is created, it is created in the first file that exists. If no files in the chain exist, then it creates the last file in the list.
3. Otherwise, `${HOME}/.kube/config` is used and no merging takes place.

#### Available Commands:

<code>current-context</code>	Displays the current-context
<code>delete-cluster</code>	Delete the specified cluster from the kubeconfig
<code>delete-context</code>	Delete the specified context from the kubeconfig
<code>get-clusters</code>	Display clusters defined in the kubeconfig
<code>get-contexts</code>	Describe one or many contexts
<code>rename-context</code>	Renames a context from the kubeconfig file.
<code>set</code>	Sets an individual value in a kubeconfig file
<code>set-cluster</code>	Sets a cluster entry in kubeconfig
<code>set-context</code>	Sets a context entry in kubeconfig
<code>set-credentials</code>	Sets a user entry in kubeconfig
<code>unset</code>	Unsets an individual value in a kubeconfig file
<code>use-context</code>	Sets the current-context in a kubeconfig file
<code>view</code>	Display merged kubeconfig settings or a specified kubeconfig file

#### Usage:

```
kubectl config SUBCOMMAND [options]
```

Use "kubectl <command> --help" for more information about a given command.

Use "kubectl options" for a list of global command-line options (applies to all commands).

```
ubuntu@ip-10-0-2-200:~$
```

Run the `kubectl config view` subcommand again to display the current client configuration.

```
ubuntu@ip-10-0-2-200:~$ kubectl config view
```

```
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: DATA+OMITTED
```

```

server: https://10.0.2.200:6443
name: kubernetes
contexts:
- context:
  cluster: kubernetes
  user: kubernetes-admin
  name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    client-certificate-data: REDACTED
    client-key-data: REDACTED
ubuntu@ip-10-0-2-200:~$

```

When you run *kubectl* commands a context is required. The context tells *kubectl* which cluster to connect to and which user to authenticate as. As you can see the values kubeadm configured means the `kubectl` command tries to reach the API server on port 6443 via our host's IP with TLS.

To view the REDACTED elements, add `--flatten`.

We can configure `kubectl` explicitly so that we can adjust our cluster settings in the future if need be. Get help on the `config set-cluster` subcommand:

```
ubuntu@ip-10-0-2-200:~$ kubectl help config set-cluster
```

Sets a cluster entry in kubeconfig.

Specifying a name that already exists will merge new fields on top of existing values for those fields.

Examples:

```
# Set only the server field on the e2e cluster entry without touching other values.
```

```
kubectl config set-cluster e2e --server=https://1.2.3.4
```

```
# Embed certificate authority data for the e2e cluster entry
```

```
kubectl config set-cluster e2e --certificate-
authority=~/.kube/e2e/kubernetes.ca.crt
```

```
# Disable cert checking for the dev cluster entry
```

```
kubectl config set-cluster e2e --insecure-skip-tls-verify=true
```

Options:

```
--embed-certs=false: embed-certs for the cluster entry in kubeconfig
```

Usage:

```
kubectl config set-cluster NAME [--server=server] [--certificate-
authority=path/to/certificate/authority]
[--insecure-skip-tls-verify=true] [options]
```

Use "kubectl options" for a list of global command-line options (applies to all commands).

```
ubuntu@ip-10-0-2-200:~$
```

`kubectl` configuration data is saved in a YAML file in your `$HOME/.kube` directory using these commands. Display the configuration file we copied in lab 2.

```
ubuntu@ip-10-0-2-200:~$ ls -la ~/.kube/
total 24
drwxrwxr-x  4 user user 4096 Dec  4 11:30 .
drwxr-xr-x 17 user user 4096 Dec  4 12:57 ..
drwxr-xr-x  3 user user 4096 Dec  4 11:30 cache
-rw-----  1 user root 5454 Dec  4 11:29 config
drwxrwxr-x  3 user user 4096 Dec  4 12:20 http-cache
ubuntu@ip-10-0-2-200:~$
```

Display the contents of the config file:

```
ubuntu@ip-10-0-2-200:~$ cat ~/.kube/config
```

```
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data:
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUN5RENDQWJDZ0F3SUJBZ0lCQURBTklna3Foa2lH0XcwQ
kFRc0ZBREFTVJNd0VRWURWUWFERXdwcmRXSmwKY201bGRHVnpNQjRYRFRFM01USXd0REU1TVRFd01sb1hEVE
kzTVRJd01qRTVNVEV3TWxvd0ZURVRNqKvHQTFRVQpBeE1LYTNWVpYSnVaWVFJY3pDQ0FTSXdEUVlKS29aSWh
2Y05BUUVCQlFBRGdnRVBBRENDQVFvQ2dnRUJBTET1CjhEdmpESFd00GxnLy9XR2pqcXl1WFQ5NDFCZ0tBNzg0
Z1c4WDZDQlJlJ0QWZQVQ1NzdaSTdHVEF3Uks0RGJMdHAkT3NhbfZFeE1KbHVzc3Z1YXAvb1BwMFlvS0hTZDByd
WRGNVZzWDdUN1NZVjJ3akl0S0JyeTB0UWpUdHJTOHo4VAp6RjVBMVVLV2lVb2RHZUhhSWtBVkr4V0NabWhRVy
80d0xoeUliVdDZChNsbnFFaStDM1M3eHNnU1hsVTJCWi8yCm8vTldHODYyVFZlMS9IZVFwNHgyNEJQUl1Q3R
tcjVuUGVaMmxGS3JvaDRJb3B3T0NUYlZlXNj10TjNidjhtYkMKemtNwjdHwU2SjBJNDNB0ENlR3liMULFZWQ5
RndmcU1RVmZnR09zN1A2RVY3L3YzVEZERDZjeXJhVXFYenNlSgpIQ2x4S29UbmPKdE4rdWJTYWxrQ0F3RUFBY
U1qTUNFd0RnWURWUjBQVVFIL0JBUURBZ0trTUE4R0ExVWRFd0VCCi93UUZNQU1CQWY4d0RRWUplb1pJaHZjTk
FRRUxUUUFEZ2dFQkFIcjU4VUwWcFlncVcrQkJEc08v0UNUei9NWkcKdXhhN1BPdEt0VUQ2V09wdWZCUXRqN2o
vR2lVaw5FNGkvTDNjVjJjNnBhYUY2TFZlQWF2VlpseWMvbERTYjZ4RQpoZFRkbFdWOG41WXFzR21pcjE3b0Zh
Uj1Ebi9HNUs5bDlzMUmhLWDAzMUsxMUwWTW1SbnJlZ0lly2p6QUY4MFZmClhiazhpZkdFdC85cnp60FovWjdib
TFh0FRpZ1JHN2ZRvmtZUwdbV2pabjFBWkRkVFlvcUNsZHVTTVRGUkF4RVQKTnBWZlB4bnBHZZBQU0l1Z2xkYk
xk0HE4VHBWOU1qUHpFN2tPendlT3ZDQ3hVbjhnLyt0ZHZadi9kZ3MxMmlScwpGL3dJUi9qdWxlakhWRUQyNEo
2UmR0aStJVURZaXZXNDQzZlNxaERpdDI4UlVoVFY5MwXUT3pyM3hRRT0KLS0tLS1FTkQgQ0VSVElGSUNBEUT
LS0tLQo=
    server: https://10.0.2.200:6443
    name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: kubernetes-admin
    name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    client-certificate-data:
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUM4akNDOWR3Z0F3SUJBZ0lJVjR10056a2RSUEV3RFFZS
```

```
ktvWklodmNOQVFFTEJRQXdGVEVUTUJFR0ExVUUKQXhNS2EzVmLaWEp1WlhSbGN6QWVGdzB4TnpFeU1EUXhPVE
V4TURKYUZ3MhHPREV5TURReE9URXhNRE5hTURReApGekFWQmd0VkJBb1REbk41YzNSbGJUcHRZWE4wWlhKek1
Sa3dGd1lEVlFRREV4QnJkV0psY201bGRHVnpMV0ZrCmJXbHVNSUlcSWpBTkJna3Foa2lHOXcwQkFRRUZBQU9D
QVE4QU1JSUJDZ0tDQVFFQXdzZ0FUQzdLa0szUlhMRWEKeWp1eG43S2pVR21LV3RwNlVTTnJnWEk5WWlkNXdwZ
WRNMct2amtSGM3WjgzYTJRUFg1SExG0EdwcEFiS2FNTgpQSLVCY205ekRKZk10Q1FKWFBhaVluQkFET3hRMT
VPL0ZiWGdLTG85ZDh1ay9Mcm5jZFUybnBhclh3dnFmdnlxCitvVdFQ3kvaTJHL0V0NkRkNlJUcmIybGk4MDl
YSXpQZW4xczQxUkZ1Yko5T01vendBa29FdkJR0UR3RC9EUUwKaWxVeGwxOFY5WVhMakVFQmFES2MxR3NS0Thq
ULZJRmVKT29TaWUxQmYveTJaMHphVHRHbzhKc3Mrb1JKdmVrbAorMGZLYU1UcGxRYUQrWmR60FcrZWRwVUZsd
1h3RTFTbi9JcDVxNVJMd1h3TVB0Qkx4Tk1GdXc5bjZzcXM0Mm9vCnB5MlNjd0lEQVFBQm95Y3dKVEFPQmd0V
hR0EJBZjhFQkFNQ0JhQXdFd1lEVlIwbEJBD3dDZ1lJS3dZQkJRUVgKQXdJd0RRWUpLb1pJaHJzTkFRRUx
EZ2dFQkFJQk1qV3lnbUtlVkl0Mk5MRkxVQnh5UDRYVExwRmZWNPaeQp2Ylc1SmcxQXNUUGZFQmcxSEhKRk8x
b2g2L21rQ05PTnZESXN3bDR2bi9UWwZlZlJBQ09pUjNHeDkrYnZlCUl4CKYydltaItbnXFvc0Nqa0ZDZGFud
jE3RzNwVFJicjhrQndrVnpKeStIL1dzWUtpMkIxcHdId3hCazJkcmLXZlIKN0ErTFd5VU5uamNnNwdiam83d1
F0bWiZTG9zMVJMMys4UU1scFlCTjMrdfZVL1p0N3l3NG9LVGtFeDNOR1RlSApXaE9kT0xaTzAxSzhWtqQ2h
hSXRITDlJa0I4ew5MS0lRcXU2UndYclJobkRrdlF1MDZTY0Q4dW4xS0VkdUs0CKVrMnFmdW0vdHprN0U2aEFB
cEhtMldalZBxcM5ju0h0Y2U1cE93Q2RqMHRJT2s5Y1NRUT0KLS0tLS1FTkQgQ0VSVElGSUNBVEUtLS0tLQo=
client-key-data:
LS0tLS1CRUdJTiBSU0EgUFJJVkFURSBLRVktLS0tLQpNSUlfCEFFJQkFBS0NBUEVBD3NnQVRDN0trSzNSWExFY
XlqdXhuN0tqVUdtS1d0cDZVU05yZ1hJOVlpZDV3VmVkc0k0Wk3Zqa21IYzdaODNhMmtQWDVITEY4R3BwQWJLYU
10UEpVQmNtOXpESmZNdENRSlhQYwZlZk1BRE94UTE1Ty8KRkxhYz0tMbZlk0HVrL0xybmNkVTJucGFyWHd2cWZ
2eXErblVXRUN5L2kyRy9FdDZEZDZSVHJiMmXpODAA5WE16UApLbjFzNDFSRnViSjlpTW96d0Frb0V2QlE5RHdE
L0RRTGlsVXhsMThwOVlyTGpFRUJhREtjMUdzUjk4a1JWSUZlCkPb1NpZTFcZi95MlowemFudEdvOEpcyctvU
kp2ZWtsKzBmS2FNVHBsUWFEEK1pkejHxK2VkcFVGbHdYd0UxU24KL0lwNXE1Ukx2WHdNUHRCTHh0TUZ1dzluNn
NxczQyb29weTJTY3dJREFRQUJBb0lCQVFDtWfudjBiNkx0NjdCTApQdzJPRHJ4ODROM2s0VUN4UUdEL1R5WjR
HZS93YTM3Vm92QzIxZEK4SS91MnlpNTM3RzdET2Q1N09VSGxIMnZGCMQxcmlhclh0YjRdajZtYlMwa3ZCcmJi
Z2VnVmlpNVczNHpYVHZTcE1rZjR2OXluSVc4RHZpZ0luRDMwWk50RGsKbktkeGxER1VsWDI4TUVuN2cxZUpEM
2lY2m1qeVAwNExFR1BkbFNyRVMxbExlNHN1R0RJUTlVTHVvanpzRkFqYgpnEudEuzIwVTZ5N2FEc0lZdVdEQ2
ZhdEx2Ym5mQy9VZSsvYnN6b1Q3dEdtRzZlYlZlX013S21UQmI0MmlaeFhECnN2RUNXYlJ2c2FpAVBmckJ2bk4
1a3M4enhRQkl1WlVYRnlzRmUxc09kb3Q0Tms5Z29MU2RZRklCaDBRV3drKy8KZnRUSGZUa0JBb0dCQU1oWGI0
cG8xZGhubUw1VnpwLWZiBUFSCEU3STRGR1ZHY0FlVXNCczUvaFI5UGhhNE5yKwp6MlFWSkTBK05ZYlJYZThPd
WxtdzRzeUpaUfN2RVBFYjlmVG1pY091clZ0MHBCSHlzYnZEY0R5Tjdhc1NLajMzCksyQzV0bGhFTFo2YmJETz
lLamU0cFNWazZWUtdVeHlQUHJyRE4rdGhvRmNddWRsTzBlcFhRQUtCQW9HQkFQamwKSENZR3NyMk1Md3Rfb1p
WQy9MS3orU2t1WU0wekRkZkhtVmNXy1VUem90bG5uVWhLN2FvYzZscjRjVEYxTkZBSQpBc0NMS3k3eLJzY05L
YlM1UFdwQ1VMTjlmSVhHZEsrSGFtT1FBN0NCsUJBK1V4SHU0eHA4WEZFNGRKR3dlUThHCmoyM1p0aXlJSjJQU
jd1VVUvNDBWTWI5cjJEdEk4b3pid09pQjdqTHpBb0dCQUxKdExPb1lkRnhlNThHY3Ft0TUKWDVRd0lpWUl3Yi
9uQ3dnUCthTm5Xekh0ekY5a2tONTZFanNRbVvk0ZDZJNDQ2Vy8vcmZnemTjcVlrMUZZbWI0agpLT3IwWjczZjF
JTHpYeXljK2E1QVliV01zaDl4RGk0aDlJQXdkRFlvZ25pLzZlNGcyM2pFK2xVGtKaDBQWkV5CkFxeFRNWHB4
ZUZ3R0VYOTRzM3dDT1FBQkFvR0FNS0ZyVjQ2MWU4eLJERTJUbUx0bTdtKzF2ak1lbk5sZDNeFkKektoSTUyV
UhLNTFRSU9EckFQTDNZMkRwbFBWR2pHQ1VVUlNnRw10Y0wrWkZnTmMweGI5QlQyQ2t3MXFjVC9PUQplUFdaa3
ZJWDFyU212SGxGakZaQ0gyaDlkajNaMlhLNXNZZjVUVWdwRWhyaHA0YnJ5NkFaZ1VKTUZJRlRtdXhoCkM4emZ
RcnNDZ1lCTExVTjRiM0xakV0Ryt5K2tFcTFmcEp4MDJhNk4zeUN5SmFyNWIZbwVjS0YvMG5nTm9PZlckCkUFk
RWE3ZzFDaVVMVjibGtZZHhxdFNiYUZEuzVERDI5MHVLOEZjWEVRdWRlcEJ1YkLDQ0ozc3ZwdndIRElHegphd
m8zc3lSd3lVdGt0MmZmYkorBHJETE8zTmg3RnJiS3N0WG5zTjgrYjVBeDBZSHhkMXRsUUE9PQotLS0tLUVORC
BSU0EgUFJJVkFURSBLRVktLS0tLQo=
ubuntu@ip-10-0-2-200:~$
```

The `kubectl config view` command will display nearly the same data, obfuscating the key data. The config file is simple and can easily be pre-generated and distributed to any client systems that require connection to a given cluster.

## 2. Test the Cluster

Now with our cluster running and `kubectl` configured lets issue some commands to test the Kubernetes cluster. The `cluster-info` subcommand can be used to test the cluster API end point and the `get nodes` command can be used to see the nodes in the cluster.

```
ubuntu@ip-10-0-2-200:~$ kubectl cluster-info
```

```
Kubernetes master is running at https://10.0.2.200:6443
KubeDNS is running at https://10.0.2.200:6443/api/v1/namespaces/kube-
system/services/kube-dns:dns/proxy
```

```
To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
ubuntu@ip-10-0-2-200:~$
```

If you are really adventurous run the suggested command for a detailed cluster overview, careful though, its a lot of information!

```
ubuntu@ip-10-0-2-200:~$ kubectl cluster-info dump |& wc -l
```

```
3413
```

```
ubuntu@ip-10-0-2-200:~$
```

To get detailed node information use the *describe node* subcommand again on the desired node name:

```
ubuntu@ip-10-0-2-200:~$ kubectl describe node ip-10-0-2-200
```

```
Name:                ip-10-0-2-200
Roles:               master
Labels:              beta.kubernetes.io/arch=amd64
                    beta.kubernetes.io/os=linux
                    kubernetes.io/hostname=ip-10-0-2-200
                    node-role.kubernetes.io/master=
Annotations:         kubeadm.alpha.kubernetes.io/cri-socket:
/var/run/dockershim.sock
                    node.alpha.kubernetes.io/ttl: 0
                    volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp:   Thu, 06 Dec 2018 14:54:36 +0000
Taints:              <none>
Unschedulable:       false
Conditions:
  Type                Status    LastHeartbeatTime         LastTransitionTime        Reason                    Message
  ----                -
NetworkUnavailable   False    Thu, 06 Dec 2018 17:01:52 +0000    Thu, 06 Dec 2018 17:01:52 +0000    WeaveIsUp                Weave pod has set this
OutOfDisk            False    Thu, 06 Dec 2018 18:27:32 +0000    Thu, 06 Dec 2018 14:54:26 +0000    KubeletHasSufficientDisk    kubelet has sufficient disk space available
MemoryPressure        False    Thu, 06 Dec 2018 18:27:32 +0000    Thu, 06 Dec 2018 14:54:26 +0000    KubeletHasSufficientMemory    kubelet has sufficient memory available
DiskPressure          False    Thu, 06 Dec 2018 18:27:32 +0000    Thu, 06 Dec 2018 14:54:26 +0000    KubeletHasNoDiskPressure     kubelet has no disk pressure
PIDPressure           False    Thu, 06 Dec 2018 18:27:32 +0000    Thu, 06 Dec 2018 14:54:26 +0000    KubeletHasSufficientPID      kubelet has sufficient PID available
```

```

Ready      True      Thu, 06 Dec 2018 18:27:32 +0000      Thu, 06 Dec 2018
17:01:59 +0000      KubeletReady      kubelet is posting ready status.
AppArmor enabled
Addresses:
  InternalIP: 10.0.2.200
  Hostname:   ip-10-0-2-200
Capacity:
  cpu:                2
  ephemeral-storage: 20263528Ki
  hugepages-2Mi:      0
  memory:             4045044Ki
  pods:              110
Allocatable:
  cpu:                2
  ephemeral-storage: 18674867374
  hugepages-2Mi:      0
  memory:             3942644Ki
  pods:              110
System Info:
  Machine ID:          9f1ba4ef4b924c148a4d816af9389de3
  System UUID:         EC2C41DA-D66D-DA59-F48A-5FBDCA03EE26
  Boot ID:             db5b7a44-fb6f-427e-a6ad-9b22812da9b6
  Kernel Version:      4.4.0-1072-aws
  OS Image:            Ubuntu 16.04.5 LTS
  Operating System:    linux
  Architecture:        amd64
  Container Runtime Version: docker://18.9.0
  Kubelet Version:     v1.13.2
  Kube-Proxy Version:  v1.13.2
Non-terminated Pods:  (8 in total)
  Namespace           Name           CPU
Requests  CPU Limits  Memory Requests  Memory Limits
-----
  kube-system         coredns-576cbf47c7-7qlht 100m (5%)
0 (0%)      70Mi (1%) 170Mi (4%)
  kube-system         coredns-576cbf47c7-qp7qk 100m (5%)
0 (0%)      70Mi (1%) 170Mi (4%)
  kube-system         etcd-ip-10-0-2-200      0 (0%)
0 (0%)      0 (0%)
  kube-system         kube-apiserver-ip-10-0-2-200 250m (12%)
0 (0%)      0 (0%)
  kube-system         kube-controller-manager-ip-10-0-2-200 200m (10%)
0 (0%)      0 (0%)
  kube-system         kube-proxy-rmxrk        0 (0%)
0 (0%)      0 (0%)
  kube-system         kube-scheduler-ip-10-0-2-200 100m (5%)
0 (0%)      0 (0%)
  kube-system         weave-net-4xsgf         20m (1%)
0 (0%)      0 (0%)
Allocated resources:
(Total limits may be over 100 percent, i.e., overcommitted.)
Resource Requests Limits
-----
cpu          770m (38%) 0 (0%)
memory      140Mi (3%) 340Mi (8%)
ubuntu@ip-10-0-2-200:~$

```



Describe provides a wealth of node information. Your report will be similar but different than the one above.

- How much memory does your node have?
- How many CPUs?
- How many pods can your node run?
- What container runtime is the `kubelet` using?
- What version of `kubelet` is your node running?

Previously we used the `version` subcommand to discover the version of the `kubectl` client but now that our config is in place we can also see the version of the cluster API Server.

```
ubuntu@ip-10-0-2-200:~$ kubectl version
```

```
Client Version: version.Info{Major:"1", Minor:"12", GitVersion:"v1.12.3",
GitCommit:"435f92c719f279a3a67808c80521ea17d5715c66", GitTreeState:"clean",
BuildDate:"2018-11-26T12:57:14Z", GoVersion:"go1.10.4", Compiler:"gc",
Platform:"linux/amd64"}
Server Version: version.Info{Major:"1", Minor:"12", GitVersion:"v1.12.3",
GitCommit:"435f92c719f279a3a67808c80521ea17d5715c66", GitTreeState:"clean",
BuildDate:"2018-11-26T12:46:57Z", GoVersion:"go1.10.4", Compiler:"gc",
Platform:"linux/amd64"}
ubuntu@ip-10-0-2-200:~$
```

If you are familiar with Golang, notice the use of the `gc` tool chain (vs `gccgo`).

### 3. Creating Applications

With our cluster running and `kubectl` configured we can try to start a simple application on the cluster. The `kubectl` command provides a `get` subcommand which can be used to get information on any one of the key Kubernetes component types: deployments, pods, replica sets, and Services. While you can type `kubectl get replicaset`, that would be fairly inhumane so `kubectl` allows you to use the abbreviation `rs` for replica sets.

If you want to save yourself even more typing. Here is tab completion without the mentioned fix.

```
ubuntu@ip-10-0-2-200:~$ kubectl get
```

```
Desktop/          .kube/          Public/
...
```

```
ubuntu@ip-10-0-2-200:~$ kubectl get
```

You can enable temporary `kubectl` bash completion with:

```
ubuntu@ip-10-0-2-200:~$ source <(kubectl completion bash)
```

```
ubuntu@ip-10-0-2-200:~$
```

And after.

```
ubuntu@ip-10-0-2-200:~$ kubectl get
certificatesigningrequest deployment networkpolicy
replicaset statefulset
...
ubuntu@ip-10-0-2-200:~$ kubectl get
```

That is much better!

In a new shell, list the currently running services, deployments, replica sets, and pods on your cluster:

```
ubuntu@ip-10-0-2-200:~$ kubectl get service,deployments,rs,pods

NAME                                TYPE             CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/kubernetes                  ClusterIP        10.96.0.1     <none>         443/TCP    3h40m
ubuntu@ip-10-0-2-200:~$
```

The only service running in our cluster is the *kubernetes* service itself. We have no deployments, replica sets, or pods yet (in our namespace). Do the same for the resources under the kube-system namespace, more on namespaces later.

```
ubuntu@ip-10-0-2-200:~$ kubectl get service,deployments,rs,pods --namespace=kube-system

NAME                                TYPE             CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/kube-dns                    ClusterIP        10.96.0.10    <none>         53/UDP,53/TCP    3h41m

NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
deployment.extensions/coredns       2/2      2              2             3h41m

NAME                                DESIRED    CURRENT    READY    AGE
replicaset.extensions/coredns-576cbf47c7  2          2          2        3h41m

NAME                                READY    STATUS    RESTARTS    AGE
pod/coredns-576cbf47c7-7qlht        1/1      Running   0            3h41m
pod/coredns-576cbf47c7-qp7qk        1/1      Running   0            3h41m
pod/etcd-ip-10-0-2-200               1/1      Running   0            3h40m
pod/kube-apiserver-ip-10-0-2-200     1/1      Running   0            3h40m
pod/kube-controller-manager-ip-10-0-2-200  1/1      Running   0            3h40m
pod/kube-proxy-rmxrk                 1/1      Running   0            3h41m
pod/kube-scheduler-ip-10-0-2-200     1/1      Running   0            3h40m
pod/weave-net-4xsgf                  2/2      Running   0            94m
ubuntu@ip-10-0-2-200:~$
```

We can view all namespaces via `--all-namespaces` (if we have permission).

To test our cluster lets run a single container pod. When configured with the Docker Engine as the container manager,

we can run any container image that Docker has preinstalled or knows how to download.

```
ubuntu@ip-10-0-2-200:~$ kubectl run my-nginx --generator=run-pod/v1 --
image=nginx:1.11 --port=80

pod/my-nginx created
ubuntu@ip-10-0-2-200:~$
```

The pod name is "my-nginx" and the image we used is "nginx", an official image pulled from Docker Hub by the Docker Engine in the background. The port switch tells Kubernetes the service port for our pod which will allow us to share the service with its users over that port (the program must actually use that port for this to work).

List the pods running on the cluster:

```
ubuntu@ip-10-0-2-200:~$ kubectl get pods

NAME          READY   STATUS    RESTARTS   AGE
my-nginx      1/1     Running   0           37s
ubuntu@ip-10-0-2-200:~$
```

This shows that our pods are deployed and up to date. It may take a bit to pull the Docker images (Ready might be 0).

You can use the `docker container ls` subcommand to display the containers running under the Docker Engine:

```
ubuntu@ip-10-0-2-200:~$ docker container ls --filter "name=nginx"

CONTAINER ID   IMAGE                                COMMAND                  CREATED
STATUS        PORTS          NAMES
01b0c8428c5c   5766334bdaa0   "nginx -g 'daemon of..." About a
minute ago    Up About a minute   k8s_my-nginx_my-
nginx_default_f2ab7342-f9e4-11e8-8781-02d9a858fbbc_0
a7b103781d85   k8s.gcr.io/pause:3.1   "/pause"                About a
minute ago    Up About a minute   k8s_POD_my-
nginx_default_f2ab7342-f9e4-11e8-8781-02d9a858fbbc_0
ubuntu@ip-10-0-2-200:~$
```

As you can see, while our `run` subcommand requested that Kubernetes run a container but 2 containers were launched at that time.

In Kubernetes, each Pod instance has an infrastructure container, which is the first container that the `kubelet` instantiates. The infrastructure container uses the image "k8s.gcr.io/pause:3.1" and acquires the pod's IP as well as a pod wide network and IPC namespace. All of the other containers in the pod then join the infrastructure container's network (`--net`) and IPC (`--ipc`) namespace allowing containers in the pod to easily communicate. The initial process ("`/pause`") that runs in the infrastructure container does nothing, its sole purpose is to act as the anchor for the pod and its shared namespaces.

You can learn more about the pause container by looking at the source and ultimately what is "`pause()`".

- <https://github.com/kubernetes/kubernetes/tree/master/build/pause>

- <https://github.com/kubernetes/kubernetes/blob/master/build/pause/pause.c>
- `man 2 pause` or <http://man7.org/linux/man-pages/man2/pause.2.html>

The Docker listing shows us 2 containers, the pod having an infrastructure container (pause) and the container we asked for (nginx).

Kubernetes gives each pod a name and reports on the pod status, the number of times the pod has been restarted and the pod's uptime. You can find the pod names embedded in the container names displayed by the `docker container ls` command:

```
ubuntu@ip-10-0-2-200:~$ docker container ls --filter "name=nginx" --format "{{.Names}}"
k8s_my-nginx_my-nginx_default_f2ab7342-f9e4-11e8-8781-02d9a858fbbc_0
k8s_POD_my-nginx_default_f2ab7342-f9e4-11e8-8781-02d9a858fbbc_0
ubuntu@ip-10-0-2-200:~$
```

Try killing the nginx container using the `docker container kill` subcommand and the ID of the underlying container based on the nginx image.

```
ubuntu@ip-10-0-2-200:~$ docker container kill \
$(docker container ls --filter "ancestor=nginx:1.11" --format {{.ID}} | head -1)
01b0c8428c5c
ubuntu@ip-10-0-2-200:~$
```

```
ubuntu@ip-10-0-2-200:~$ docker container ls --filter "name=nginx"

CONTAINER ID        IMAGE                                     COMMAND                  CREATED
STATUS             PORTS                                     NAMES                    ago
9a1a44e5d0f1        5766334bdaa0                            "nginx -g 'daemon of..." 9 seconds ago
Up 9 seconds                                     k8s_my-nginx_my-nginx_default_f2ab7342-
f9e4-11e8-8781-02d9a858fbbc_1
a7b103781d85        k8s.gcr.io/pause:3.1                    "/pause"                 4 minutes ago
Up 4 minutes                                     k8s_POD_my-nginx_default_f2ab7342-f9e4-
11e8-8781-02d9a858fbbc_0
ubuntu@ip-10-0-2-200:~$
```

We can tell by the created time we have a new container. If you were fast enough, you may have seen the previous container exited. Docker terminates the container specified but Kubernetes has no knowledge of this action. When the Kubelet process, responsible for the pods assigned to this node, sees the missing container, it simply reruns the nginx image.

After some time, if you run the previous command with the `-a` flag, we can see the previous killed container and the newly created one.

```
ubuntu@ip-10-0-2-200:~$ docker container ls -a --filter "name=nginx"
```

CONTAINER ID	IMAGE	COMMAND	CREATED
9a1a44e5d0f1	5766334bdaa0	"nginx -g 'daemon of...'"	About a minute ago
01b0c8428c5c	5766334bdaa0	"nginx -g 'daemon of...'"	5 minutes ago
nginx_my-nginx_default_f2ab7342-f9e4-11e8-8781-02d9a858fbbc_1			
nginx_default_f2ab7342-f9e4-11e8-8781-02d9a858fbbc_0			
a7b103781d85	k8s.gcr.io/pause:3.1	"/pause"	5 minutes ago
nginx_my-nginx_my-nginx_default_f2ab7342-f9e4-11e8-8781-02d9a858fbbc_0			
ubuntu@ip-10-0-2-200:~\$			

Notice that we killed container `01b0c8428c5c` in the example but the new container `9a1a44e5d0f1` was created to replace it. Kubernetes does not "resurrect" containers that have failed. This is important because the container's state may be the reason it failed. Rather, Kubernetes runs a fresh copy of the original image, ensuring the container has a clean new internal state (cattle not pets!).

## 4. Create a Service

In modern software engineering terms, a service is an encapsulated set of functionality made available to consumers through an API. The problem with our nginx application at present is that when containers die new ones are created. The fact that there are multiple containers and that containers come and go makes using the app difficult.

To simplify things Kubernetes makes it possible for us to expose our pods as a Service. The `kubectl expose` command does this.

Expose the my-nginx pod replica set as a service:

```
ubuntu@ip-10-0-2-200:~$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
my-nginx	1/1	Running	1	3m38s

```
ubuntu@ip-10-0-2-200:~$
```

```
ubuntu@ip-10-0-2-200:~$ kubectl expose $(kubectl get pod -o=name) --port=80
```

```
service/my-nginx exposed
ubuntu@ip-10-0-2-200:~$
```

This causes Kubernetes to create a conceptual Service for our pods, exposing the set of pods as a single endpoint for users. Use the `get services` subcommand to display your service.

```
ubuntu@ip-10-0-2-200:~$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4h25m
my-nginx-76bcc6d46	ClusterIP	10.100.43.206	<none>	80/TCP	15s

```
ubuntu@ip-10-0-2-200:~$
```

Kubernetes has given our service a virtual IP (VIP) address and it will now distribute client connections across the pods running my-nginx.

To test the Service try curling it:

```
ubuntu@ip-10-0-2-200:~$ NX_CIP=$(kubectl get services -o=custom-  
columns=NAME:.spec.clusterIP,NAME:.metadata.name \  
| grep nginx | awk '{print $1}') && echo $NX_CIP  
  
10.100.43.206  
  
ubuntu@ip-10-0-2-200:~$ curl -I $NX_CIP  
  
HTTP/1.1 200 OK  
Server: nginx/1.11.13  
Date: Wed, 28 Mar 2018 22:32:41 GMT  
Content-Type: text/html  
Content-Length: 612  
Last-Modified: Tue, 04 Apr 2017 15:01:57 GMT  
Connection: keep-alive  
ETag: "58e3b565-264"  
Accept-Ranges: bytes  
ubuntu@ip-10-0-2-200:~$
```

Success!

## 5. Pod exec

While Kubernetes delegates all of the direct container operations to the container manager (usually Docker) it does pass through some useful container features.

For example, imagine you need to discover the distro of one of your pods' containers. You can use the `kubectl exec` subcommand to run arbitrary commands within a pod.

Try listing the running pods and then executing the `cat /etc/os-release` command within one of your pods.

```
ubuntu@ip-10-0-2-200:~$ kubectl get pods  
  
NAME          READY   STATUS    RESTARTS   AGE  
my-nginx      1/1     Running   1           8m38s  
  
ubuntu@ip-10-0-2-200:~$ kubectl exec my-nginx cat /etc/os-release  
  
PRETTY_NAME="Debian GNU/Linux 8 (jessie)"  
NAME="Debian GNU/Linux"  
VERSION_ID="8"  
VERSION="8 (jessie)"  
ID=debian  
HOME_URL="http://www.debian.org/"  
SUPPORT_URL="http://www.debian.org/support"  
BUG_REPORT_URL="https://bugs.debian.org/"
```

```
ubuntu@ip-10-0-2-200:~$
```

Running `cat /etc/os-release` via `kubectl exec` produces the information we needed. The `exec` subcommand chooses the first container within the pod to execute the command.

If you would like to execute the command within a specific container you can use the `-c` switch. The *describe pod* command will give you a list of the containers within the pod. We can also retrieve JSON output and filter for it. Our current pod has only one container but we can still test the command.

Try it:

```
ubuntu@ip-10-0-2-200:~$ kubectl get pod my-nginx -o json | jq
".spec.containers[].name" -r

my-nginx
ubuntu@ip-10-0-2-200:~$
```

Use the `-c` switch to display the `os-release` file in the my-nginx container in the pod:

```
ubuntu@ip-10-0-2-200:~$ kubectl exec -c my-nginx my-nginx cat /etc/os-release

PRETTY_NAME="Debian GNU/Linux 8 (jessie)"
NAME="Debian GNU/Linux"
VERSION_ID="8"
VERSION="8 (jessie)"
ID=debian
HOME_URL="http://www.debian.org/"
SUPPORT_URL="http://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"
ubuntu@ip-10-0-2-200:~$
```

## 6. System Logs

Each of the services composing our Kubernetes cluster emits a log file. In the current configuration, the `kubelet` log is controlled by systemd.

You can use the `journalctl` command to tail ( `-n` ) the output for the kubelet service unit ( `-u` )

```
ubuntu@ip-10-0-2-200:~$ journalctl -n 400 --no-pager -u kubelet.service | grep -v
"no observation"

-- Logs begin at Wed 2018-03-28 20:50:52 UTC, end at Thu 2018-03-29 01:39:50 UTC.
--
Mar 28 22:11:54 ip-10-0-2-200 kubelet[6783]: Flag --client-ca-file has been
deprecated, This parameter should be set via the config file specified by the
Kubelet's --config flag. See https://kubernetes.io/docs/tasks/administer-
cluster/kubelet-config-file/ for more information.
Mar 28 22:11:54 ip-10-0-2-200 kubelet[6783]: Flag --cadvisor-port has been
deprecated, The default will change to 0 (disabled) in 1.12, and the cadvisor
```

```

port will be removed entirely in 1.13
Mar 28 22:11:54 ip-10-0-2-200 kubelet[6783]: I0328 22:11:54.739979    6783
feature_gate.go:226] feature gates: &{{}} map[]
Mar 28 22:11:54 ip-10-0-2-200 kubelet[6783]: F0328 22:11:54.740056    6783
server.go:218] unable to load client CA file /etc/kubernetes/pki/ca.crt: open
/etc/kubernetes/pki/ca.crt: no such file or directory
Mar 28 22:11:54 ip-10-0-2-200 systemd[1]: kubelet.service: Main process exited,
code=exited, status=255/n/a
Mar 28 22:11:54 ip-10-0-2-200 systemd[1]: kubelet.service: Unit entered failed
state.
Mar 28 22:11:54 ip-10-0-2-200 systemd[1]: kubelet.service: Failed with result
'exit-code'.
Mar 28 22:12:04 ip-10-0-2-200 systemd[1]: kubelet.service: Service hold-off time
over, scheduling restart.
Mar 28 22:12:04 ip-10-0-2-200 systemd[1]: Stopped kubelet: The Kubernetes Node
Agent.
Mar 28 22:12:04 ip-10-0-2-200 systemd[1]: Started kubelet: The Kubernetes Node
Agent.
Mar 28 22:12:04 ip-10-0-2-200 kubelet[6791]: Flag --pod-manifest-path has been
deprecated, This parameter should be set via the config file specified by the
Kubelet's --config flag. See https://kubernetes.io/docs/tasks/administer-
cluster/kubelet-config-file/ for more information.
Mar 28 22:12:04 ip-10-0-2-200 kubelet[6791]: Flag --allow-privileged has been
deprecated, will be removed in a future version

...

ubuntu@ip-10-0-2-200:~$

```

The rest of our services are running as containers.

We can use the `kubectl logs` command to display log output from our pods. Remember that Kubernetes system services run within the `kube-system` namespace by convention.

List the pods in the `kube-system` namespace:

```

ubuntu@ip-10-0-2-200:~$ kubectl get pods --namespace=kube-system

NAME                                READY    STATUS    RESTARTS   AGE
etcd-ip-10-0-2-200                 1/1      Running   0           3h
kube-apiserver-ip-10-0-2-200        1/1      Running   0           3h
kube-controller-manager-ip-10-0-2-200 1/1      Running   0           3h
kube-dns-86f4d74b45-dqfs5          3/3      Running   0           3h
kube-proxy-xw6nh                   1/1      Running   0           3h
kube-scheduler-ip-10-0-2-200        1/1      Running   0           3h
weave-net-rmsx8                    2/2      Running   0           3h
ubuntu@ip-10-0-2-200:~$

```

Now display the last 10 lines from the API service:

```

ubuntu@ip-10-0-2-200:~$ kubectl logs --namespace=kube-system --tail=10 kube-
apiserver-ip-10-0-2-200

```



```

I0124 21:20:21.743616      1 storage_rbac.go:276] created
rolebinding.rbac.authorization.k8s.io/system:controller:bootstrap-signer in kube-
system
I0124 21:20:21.785787      1 storage_rbac.go:276] created
rolebinding.rbac.authorization.k8s.io/system:controller:cloud-provider in kube-
system
I0124 21:20:21.825880      1 storage_rbac.go:276] created
rolebinding.rbac.authorization.k8s.io/system:controller:token-cleaner in kube-
system
I0124 21:20:21.866111      1 storage_rbac.go:276] created
rolebinding.rbac.authorization.k8s.io/system:controller:bootstrap-signer in kube-
public
I0124 21:20:22.119230      1 controller.go:608] quota admission added evaluator
for: serviceaccounts
I0124 21:20:23.152183      1 controller.go:608] quota admission added evaluator
for: deployments.apps
I0124 21:20:23.189041      1 controller.go:608] quota admission added evaluator
for: daemonsets.apps
I0124 21:20:28.565046      1 controller.go:608] quota admission added evaluator
for: replicaset.apps
I0124 21:20:28.669494      1 controller.go:608] quota admission added evaluator
for: controllerrevisions.apps
I0124 21:30:43.442529      1 controller.go:608] quota admission added evaluator
for: daemonsets.extensions
ubuntu@ip-10-0-2-200:~$

```

Each Kubernetes service has its own log verbosity and each can be tuned. You can learn much by tracking the operations involved in starting a deployment.

Create a new single pod deployment with a descriptive name and then grep its activity in the logs.

```

ubuntu@ip-10-0-2-200:~$ kubectl run --generator=run-pod/v1 mylogtracker --image
nginx:1.11

pod/mylogtracker created
ubuntu@ip-10-0-2-200:~$

```

Again list the k8s system services, from here we can pick which logs to search for our new pod.

```

ubuntu@ip-10-0-2-200:~$ kubectl get pod --namespace=kube-system

```

NAME	READY	STATUS	RESTARTS	AGE
etcd-ip-10-0-2-200	1/1	Running	0	3h
kube-apiserver-ip-10-0-2-200	1/1	Running	0	3h
kube-controller-manager-ip-10-0-2-200	1/1	Running	0	3h
kube-dns-86f4d74b45-dqfs5	3/3	Running	0	3h
kube-proxy-xw6nh	1/1	Running	0	3h
kube-scheduler-ip-10-0-2-200	1/1	Running	0	3h
weave-net-rmsx8	2/2	Running	0	3h

```

ubuntu@ip-10-0-2-200:~$

```

Try the controller manager server first:

```
ubuntu@ip-10-0-2-200:~$ kubectl logs --namespace=kube-system kube-controller-manager-ip-10-0-2-200 | grep mylogtracker
```

```
ubuntu@ip-10-0-2-200:~$
```

Controller manager only deals with replicated pods (ones using a controller); there won't be anything here for us.

Now take a look at the kubelet log:

```
ubuntu@ip-10-0-2-200:~$ journalctl -u kubelet.service | grep mylogtracker
```

```
Dec 04 13:28:35 ip-10-0-2-200 kubelet[74364]: I1204 13:28:35.182932 74364 reconciler.go:212] operationExecutor.VerifyControllerAttachedVolume started for volume "default-token-2kmhb" (UniqueName: "kubernetes.io/secret/1584a8f2-d93a-11e7-a277-000c29ae8ddc-default-token-2kmhb") pod "mylogtracker-6ff4ff6fd5-k5qrq" (UID: "1584a8f2-d93a-11e7-a277-000c29ae8ddc")
ubuntu@ip-10-0-2-200:~$
```

Kubelet only reports information about our pod's volume.

You can also view the events taking place within the Kubernetes cluster itself using the events resource type.

Try getting events with `kubectl`:

```
ubuntu@ip-10-0-2-200:~$ kubectl get events
```

LAST SEEN	TYPE	REASON	KIND	MESSAGE
11m	Normal	Scheduled	Pod	Successfully assigned
default/my-nginx to ubuntu				
11m	Normal	Pulling	Pod	pulling image "nginx:1.11"
10m	Normal	Pulled	Pod	Successfully pulled image
"nginx:1.11"				
9m13s	Normal	Created	Pod	Created container
9m13s	Normal	Started	Pod	Started container
9m13s	Normal	Pulled	Pod	Container image
"nginx:1.11" already present on machine				
81s	Normal	Scheduled	Pod	Successfully assigned
default/mylogtracker to ubuntu				
81s	Normal	Pulled	Pod	Container image
"nginx:1.11" already present on machine				
81s	Normal	Created	Pod	Created container
80s	Normal	Started	Pod	Started container
31m	Normal	Starting	Node	Starting kubelet.
31m	Normal	NodeHasSufficientMemory	Node	Node ubuntu status is now:
NodeHasSufficientMemory				
31m	Normal	NodeHasNoDiskPressure	Node	Node ubuntu status is now:
NodeHasNoDiskPressure				
31m	Normal	NodeHasSufficientPID	Node	Node ubuntu status is now:
NodeHasSufficientPID				
31m	Normal	NodeAllocatableEnforced	Node	Updated Node Allocatable
limit across pods				
31m	Normal	RegisteredNode	Node	Node ubuntu event:

```
Registered Node ubuntu in Controller
31m      Normal    Starting
20m      Normal    NodeReady
NodeReady
ubuntu@ip-10-0-2-200:~$
```

```
Node    Starting kube-proxy.
Node    Node ubuntu status is now:
```

While your events will be different you can see the value of the cluster event log. You can display event data associated with a given resource by supplying its name. You can also control the output format.

For example to make the data machine readable you could output it in JSON:

```
ubuntu@ip-10-0-2-200:~$ kubectl get -o json events | tail
      },
      "type": "Normal"
    }
  ],
  "kind": "List",
  "metadata": {
    "resourceVersion": "",
    "selfLink": ""
  }
}
ubuntu@ip-10-0-2-200:~$
```

## 7. Cleaning Up

Now that we have given our new cluster a good test we can clean up by deleting the service and deployments we have created. The `kubectl delete` subcommand allows you to delete objects you have created in the cluster.

To begin, delete the *my-nginx* Service:

```
ubuntu@ip-10-0-2-200:~$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	5h
my-nginx-87464966f	ClusterIP	10.102.46.131	<none>	80/TCP	8m

```
ubuntu@ip-10-0-2-200:~$
```

```
ubuntu@ip-10-0-2-200:~$ kubectl delete service my-nginx
```

```
service "my-nginx" deleted
ubuntu@ip-10-0-2-200:~$
```

```
ubuntu@ip-10-0-2-200:~$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
------	------	------------	-------------	---------	-----

```
svc/kubernetes    ClusterIP    10.96.0.1    <none>        443/TCP    2h
ubuntu@ip-10-0-2-200:~$
```

Do not delete the kubernetes service.

Next we can delete the deployments (which in turn removes the associated replica set).

```
ubuntu@ip-10-0-2-200:~$ kubectl get pod

NAME            READY   STATUS    RESTARTS   AGE
my-nginx        1/1     Running   1           17m
mylogtracker    1/1     Running   0           5m31s
ubuntu@ip-10-0-2-200:~$
```

```
ubuntu@ip-10-0-2-200:~$ kubectl delete pod my-nginx

pod "my-nginx" deleted
ubuntu@ip-10-0-2-200:~$
```

```
ubuntu@ip-10-0-2-200:~$ kubectl delete pod mylogtracker

pod "mylogtracker" deleted
ubuntu@ip-10-0-2-200:~$
```

```
ubuntu@ip-10-0-2-200:~$ kubectl get pods

No resources found.
ubuntu@ip-10-0-2-200:~$
```

You Kubernetes cluster should now be cleaned up and ready for the next lab:

```
ubuntu@ip-10-0-2-200:~$ kubectl get services,deployments,replicasets,pods

NAME                                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
service/kubernetes                 ClusterIP     10.96.0.1     <none>         443/TCP    34m
user@ubuntu:~$
```

Be sure to leave the `service/kubernetes` service!

When we delete a deployment, Kubernetes ensures that all pods controlled by that replica set are also deleted. If you have trouble with deleting resources, the order matters. Deployments watch replica sets, replica sets watch pods; start with deployments. Services fall outside of that restart logic, and can be deleted at any point.

Congratulations, you have completed the lab!

*Copyright (c) 2013-2019 RX-M LLC, Cloud Native Consulting, all rights reserved. Licensed for use only by students in Safewrd Ventures class 18-04-19 (Max 10 students)*