

Kubernetes

Lab 3 – Working with Pods

In this lab we will explore the nature of Kubernetes pods and how to work with them.

In Kubernetes, pods are the smallest deployable units that can be created, scheduled, and managed. A pod corresponds to a collocated group of containers running with a shared context. Within that context, the applications may also have individual cgroup isolations applied. A pod models an application-specific "logical host" in a containerized environment. It may contain one or more applications which are relatively tightly coupled — in a pre-container world, they would have executed on the same physical or virtual host.

The context of the pod can be defined as the conjunction of several Linux namespaces:

- PID - applications within the pod can see each other's processes (as of Docker 1.12)
- Network - applications within the pod have access to the same IP and port space
- IPC - applications within the pod can use SystemV IPC or POSIX message queues to communicate
- UTS - applications within the pod share a hostname

Applications within a pod can also have access to shared volumes, which are defined at the pod level and made available in each application's file system. Additionally, a pod may define top-level cgroup isolations which form an outer bound to any individual isolation applied to constituent containers.

Like individual application containers, pods are considered to be relatively ephemeral rather than durable entities. Pods are scheduled to nodes and remain there until termination (according to restart policy) or deletion. When a node dies, the s scheduled to that node are deleted. Specific pods are never moved to new nodes; instead, they must be replaced by running fresh copies of the images on the new node.

As a first step in our exploration we will create a simple single container pod.

1. A Simple Pod

To begin our exploration, we'll create a basic Kubernetes pod from the command line. The easiest way to run a pod is using the `kubectl run` command. Try creating a simple Apache Web Server pod using the `kubectl run` subcommand as follows.

```
ubuntu@ip-10-0-2-200:~$ kubectl run apache --generator=run-pod/v1 --  
image=httpd:2.2  
  
pod/apache created  
ubuntu@ip-10-0-2-200:~$
```

Now view the pod:

```
ubuntu@ip-10-0-2-200:~$ kubectl get pod
```

| NAME | READY | STATUS | RESTARTS | AGE |
|--------|-------|---------|----------|-----|
| apache | 1/1 | Running | 0 | 1m |

```
ubuntu@ip-10-0-2-200:~$
```

What happened here?

The run command takes a name, an image and an API object as parameters and it generates pod template including the image you specified.

The `run` subcommand syntax is as follows:

```
ubuntu@ip-10-0-2-200:~$ kubectl run -h | grep COMMAND
```

```

  kubectl run NAME --image=image [--env="key=value"] [--port=port] [--
  replicas=replicas] [--dry-run=bool] [--overrides=inline-json] [--command] --
  [COMMAND] [args...] [options]
ubuntu@ip-10-0-2-200:~$
```

The `--env` switch sets environment variables (just like the `docker run -e` switch,) the `--port` switch exposes ports for service mapping, the `--replicas` switch sets the number of instances of the pod you would like the cluster to maintain and the `--dry-run` switch (if set to true) allows you to submit the command without executing it to test the parameters.

It doesn't include anything about the `--generator=` flag, let's try again:

```
ubuntu@ip-10-0-2-200:~$ kubectl run -h | grep generator
```

```

  --generator='': The name of the API generator to use, see
  http://kubernetes.io/docs/user-guide/kubectl-conventions/#generators for a list.
  --service-generator='service/v2': The name of the generator to use for
  creating a service. Only used if --expose is true
ubuntu@ip-10-0-2-200:~$
```

Taking a look at the link provided by the help text reveals a number of different generators:

| Resource | kubectl command |
|----------------------------|---|
| Pod | <code>kubectl run --generator=run-pod/v1</code> |
| Replication controller | <code>kubectl run --generator=run/v1</code> |
| Deployment | <code>kubectl run --generator=extensions/v1beta1</code> |
| -for an endpoint (default) | <code>kubectl run --generator=deployment/v1beta1</code> |
| Deployment | <code>kubectl run --generator=apps/v1beta1</code> |
| | |

| | |
|--------------------------------|---|
| -for an endpoint (recommended) | kubectl run --generator=deployment/apps.v1beta1 |
| Job | kubectl run --generator=job/v1 |
| CronJob | kubectl run --generator=batch/v1beta1 |
| -for an endpoint (default) | kubectl run --generator=cronjob/v1beta1 |
| CronJob | kubectl run --generator=batch/v2alpha1 |
| -for an endpoint (deprecated) | kubectl run --generator=cronjob/v2alpha1 |

We can use `run` to create Pods or Controllers (RCs, Deployments, Jobs, CronJobs); more on Controllers later.

To get more information about our pod use the `kubectl describe` subcommand:

```
ubuntu@ip-10-0-2-200:~$ kubectl describe pod apache

Name:          apache
Namespace:     default
Priority:       0
PriorityClassName:  <none>
Node:          ip-10-0-2-200/10.0.2.200
Start Time:    Thu, 06 Dec 2018 19:34:57 +0000
Labels:        run=apache
Annotations:    <none>
Status:        Running
IP:            10.32.0.4
Containers:
  apache:
    Container ID:  docker://f57161648fbfb6b30750d6f00a497554e9f743e601e67ea969504269e52a0606
    Image:         httpd:2.2
    Image ID:      docker-
    pullable://httpd@sha256:9784d70c8ea466fabd52b0bc8cde84980324f9612380d22fbad2151df9a430eb
    Port:          <none>
    Host Port:     <none>
    State:         Running
      Started:     Thu, 06 Dec 2018 19:34:58 +0000
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-pdn99 (ro)
Conditions:
  Type            Status
  Initialized      True
  Ready           True
  ContainersReady True
  PodScheduled     True
Volumes:
```

```

default-token-pdn99:
  Type:          Secret (a volume populated by a Secret)
  SecretName:    default-token-pdn99
  Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s

Events:
  Type       Reason      Age   From                      Message
  ----       -
Normal      Scheduled   8m37s default-scheduler         Successfully assigned
default/apache to ip-10-0-2-200
Normal      Pulled      8m37s kubelet, ip-10-0-2-200    Container image "httpd:2.2"
already present on machine
Normal      Created     8m37s kubelet, ip-10-0-2-200    Created container
Normal      Started     8m36s kubelet, ip-10-0-2-200    Started container
ubuntu@ip-10-0-2-200:~$

```

Read through the *Events* reported for the pod.

You can see that Kubernetes used the Docker Engine to pull, create, and start the httpd image requested. You can also see which part of Kubernetes caused the event. For example the scheduler assigned the pod to node Ubuntu and then the Kubelet on node Ubuntu starts the container.

Use the `docker container ls` subcommand to examine your containers directly on the Docker Engine.

```

ubuntu@ip-10-0-2-200:~$ docker container ls -f "name=apache"

CONTAINER ID        IMAGE               COMMAND             CREATED
STATUS              PORTS              NAMES
f57161648fbf        e06c3dbbfe23       "httpd-foreground"  9 minutes ago
Up 9 minutes        k8s_apache_apache_default_034884dd-f98e-
11e8-8781-02d9a858fbbc_0
a615674ea9d0        k8s.gcr.io/pause:3.1  "/pause"           9 minutes ago
Up 9 minutes        k8s_POD_apache_default_034884dd-f98e-
11e8-8781-02d9a858fbbc_0
ubuntu@ip-10-0-2-200:~$

```

Kubernetes incorporates the pod name into the name of each container running in the pod.

Use the `docker container inspect` subcommand to examine the container details for your httpd container:

```

ubuntu@ip-10-0-2-200:~$ docker container inspect $(docker container ls -f
"name=apache" --format '{{.ID}}') | less
[
  {
    "Id": "f57161648fbfb6b30750d6f00a497554e9f743e601e67ea969504269e52a0606",
    "Created": "2018-12-06T19:34:57.803581232Z",
    "Path": "httpd-foreground",
    "Args": [],
    "State": {
      "Status": "running",

```

```

        "Running": true,
        "Paused": false,
        "Restarting": false,
        "OOMKilled": false,
        "Dead": false,
        "Pid": 29340,
        "ExitCode": 0,
        "Error": "",
        "StartedAt": "2018-12-06T19:34:58.05454054Z",
        "FinishedAt": "0001-01-01T00:00:00Z"
    },
    "Image":
"sha256:e06c3dbbfe239c6fca50b6ab6935b3122930fa2eea2136979e5b46ad77ecb685",
    "ResolvConfPath":
"/var/lib/docker/containers/a615674ea9d0bb051056b581326667126eba6bbd498a5c19c8b9d
ab6e9ffb3a3/resolv.conf",
    "HostnamePath":
"/var/lib/docker/containers/a615674ea9d0bb051056b581326667126eba6bbd498a5c19c8b9d
ab6e9ffb3a3/hostname",
    "HostsPath": "/var/lib/kubelet/pods/034884dd-f98e-11e8-8781-
02d9a858fbbc/etc-hosts",
    "LogPath":
"/var/lib/docker/containers/f57161648fbfb6b30750d6f00a497554e9f743e601e67ea969504
269e52a0606/f57161648fbfb6b30750d6f00a497554e9f743e601e67ea969504269e52a0606-
json.log",
    "Name": "/k8s_apache_apache_default_034884dd-f98e-11e8-8781-
02d9a858fbbc_0",
    "RestartCount": 0,
    "Driver": "overlay2",
    "Platform": "linux",
    "MountLabel": "",
    "ProcessLabel": "",
    "AppArmorProfile": "docker-default",
    "ExecIDs": null,
    "HostConfig": {
        "Binds": [
            "/var/lib/kubelet/pods/034884dd-f98e-11e8-8781-
02d9a858fbbc/volumes/kubernetes.io~secret/default-token-
pdn99:/var/run/secrets/kubernetes.io/serviceaccount:ro",
            "/var/lib/kubelet/pods/034884dd-f98e-11e8-8781-02d9a858fbbc/etc-
hosts:/etc/hosts",
            "/var/lib/kubelet/pods/034884dd-f98e-11e8-8781-
02d9a858fbbc/containers/apache/88746e33:/dev/termination-log"
        ],
        "ContainerIDFile": "",
        "LogConfig": {
            "Type": "json-file",
            "Config": {}
        }
    },
    ...
}
ubuntu@ip-10-0-2-200:~$

```

Notice that Kubernetes injects a standard set of environment variables into the containers of a pod providing the location of the cluster API service.

```

ubuntu@ip-10-0-2-200:~$ docker container inspect $(docker container ls -f
"name=apache" --format '{{.ID}}') \
| jq -r '.[0].Config.Env[]'

KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_PORT=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP_PROTO=tcp
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
PATH=/usr/local/apache2/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sb
in:/bin
HTTPD_PREFIX=/usr/local/apache2
HTTPD_VERSION=2.2.34
HTTPD_SHA256=e53183d5dfac5740d768b4c9bea193b1099f4b06b57e5f28d7caaf9ea7498160
HTTPD_PATCHES=CVE-2017-9798-patch-2.2.patch
42c610f8a8f8d4d08664db6d9857120c2c252c9b388d56f238718854e6013e46 2.2.x-mod_proxy-
without-APR_HAS_THREADS.patch
beb66a79a239f7e898311c5ed6a38c070c641ec56706a295b7e5caf3c55a7296
APACHE_DIST_URLS=https://www.apache.org/dyn/closer.cgi?action=download&filename=
https://www-us.apache.org/dist/ https://www.apache.org/dist/
https://archive.apache.org/dist/
ubuntu@ip-10-0-2-200:~$

```

Kubernetes has also added several labels to the container, for example you can see that the pod is running in the default namespace.

```

ubuntu@ip-10-0-2-200:~$ docker container inspect $(docker container ls -f
"name=apache" --format '{{.ID}}') \
| jq -r '.[0].Config.Labels'

{
  "annotation.io.kubernetes.container.hash": "ed835021",
  "annotation.io.kubernetes.container.restartCount": "0",
  "annotation.io.kubernetes.container.terminationMessagePath": "/dev/termination-
log",
  "annotation.io.kubernetes.container.terminationMessagePolicy": "File",
  "annotation.io.kubernetes.pod.terminationGracePeriod": "30",
  "io.kubernetes.container.logpath": "/var/log/pods/034884dd-f98e-11e8-8781-
02d9a858fbbc/apache/0.log",
  "io.kubernetes.container.name": "apache",
  "io.kubernetes.docker.type": "container",
  "io.kubernetes.pod.name": "apache",
  "io.kubernetes.pod.namespace": "default",
  "io.kubernetes.pod.uid": "034884dd-f98e-11e8-8781-02d9a858fbbc",
  "io.kubernetes.sandbox.id":
"a615674ea9d0bb051056b581326667126eba6bbd498a5c19c8b9dab6e9ffb3a3"
}
ubuntu@ip-10-0-2-200:~$

```

curl the IP address of the web server to ensure that you can reach the running Apache web server. Don't forget - complete.

```
ubuntu@ip-10-0-2-200:~$ PIP=$(kubectl get pod -o jsonpath='{.items[*].status.podIP}') && echo $PIP
```

```
10.32.0.4
ubuntu@ip-10-0-2-200:~$
```

```
ubuntu@ip-10-0-2-200:~$ curl -I $PIP
```

```
HTTP/1.1 200 OK
Date: Thu, 29 Mar 2018 05:19:31 GMT
Server: Apache/2.2.34 (Unix) mod_ssl/2.2.34 OpenSSL/1.0.1t DAV/2
Last-Modified: Sat, 20 Nov 2004 20:16:24 GMT
ETag: "82a0d-2c-3e9564c23b600"
Accept-Ranges: bytes
Content-Length: 44
Content-Type: text/html
ubuntu@ip-10-0-2-200:~$
```

- How can you discover the address of the Apache web server running inside the container?
 - try

```
docker container inspect $(docker container ls -f "ancestor=httpd:2.2" --format '{{.ID}}') |grep IPAddress'
```
- Why can't you find the IP address in the container inspect data?
- What is the value of the Apache container's HostConfig.NetworkMode setting in the inspect data?
 - hint:

```
docker container inspect $(docker container ls -f "ancestor=httpd:2.2" --format '{{.ID}}') -f '{{.HostConfig.NetworkMode}}'
```
- What container does this reference?
 - ```
docker container ls -f "ID=<insert_ID_returned_by_the_previous_command_here>"
```
- What other keys in the apache container inspect has the same value and why?

Because pods house running processes on nodes in the cluster, it is important to allow those processes to gracefully terminate when they are no longer needed. In Kubernetes, users can request deletion and discover when processes terminate. When a user requests deletion of a pod, Kubernetes sends the appropriate termination signal to each container in the pod (either SIGTERM or the container defined stop signal). Kubernetes then waits for a grace period after which, if the pod has not shutdown, the pod is forcefully killed with SIGKILL and the pod is then deleted from the API server. If the Kubelet or the container manager is restarted while waiting for processes to terminate, the termination will be retried with the full grace period.

While `kubectl run` is handy for quickly starting a single image based pod it is not the most flexible or repeatable way to create pods. Next we'll take a look at a more useful way of starting pods, from a configuration file.

## 2. Pod Config Files

Kubernetes supports declarative YAML or JSON configuration files. Often times config files are preferable to imperative commands, since they can be checked into version control and changes to the files can be code reviewed, producing a more robust, reliable and CI/CD friendly system. They can also save a lot of typing if you would like to

deploy complex pods or entire applications.

Let's try running an nginx container in a pod but this time we'll create the pod using a YAML configuration file. Create the following config file with your favorite editor (as long as it is vim) in a new "pods" working directory.

```
ubuntu@ip-10-0-2-200:~$ mkdir pods
ubuntu@ip-10-0-2-200:~$
```

```
ubuntu@ip-10-0-2-200:~$ cd pods/
ubuntu@ip-10-0-2-200:~/pods$
```

```
ubuntu@ip-10-0-2-200:~/pods$ vi nginxpod.yaml
ubuntu@ip-10-0-2-200:~/pods$ cat nginxpod.yaml

apiVersion: v1
kind: Pod
metadata:
 name: nginxpod
spec:
 containers:
 - name: nginx
 image: nginx:1.11
 ports:
 - containerPort: 80
ubuntu@ip-10-0-2-200:~/pods$
```

The key "kind" tells Kubernetes we wish to create a pod. The "metadata" section allows us to define a name for the pod and to apply any other labels we might deem useful.

The "spec" section defines the containers we wish to run in our pod. In our case we will run just a single container based on the nginx image. The "ports" key allows us to share the ports the pod will be using with the orchestration layer. More on this later.

To have Kubernetes create your new pod you can use the `kubectl create` subcommand. The `create` subcommand will accept a config file via stdin or you can load the config from a file with the `-f` switch (more common). Try the following:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl create -f nginxpod.yaml

pod/nginxpod created
ubuntu@ip-10-0-2-200:~/pods$
```

Now list the pods on your cluster:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl get pods
```



| NAME     | READY | STATUS  | RESTARTS | AGE |
|----------|-------|---------|----------|-----|
| nginxpod | 1/1   | Running | 0        | 15s |

ubuntu@ip-10-0-2-200:~/pods\$

Describe your pod:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl describe pod nginxpod
```

```
Name: nginxpod
Namespace: default
Priority: 0
PriorityClassName: <none>
Node: ip-10-0-2-200/10.0.2.200
Start Time: Thu, 06 Dec 2018 21:25:40 +0000
Labels: <none>
Annotations: <none>
Status: Running
IP: 10.32.0.4
Containers:
 nginx:
 Container ID: docker://b4469f39efed8d54bb3e33398d9bf0f44ff6b439496a8a26d766e168be03b618
 Image: nginx:1.11
 Image ID: docker-
 pullable: //nginx@sha256:e6693c20186f837fc393390135d8a598a96a833917917789d63766cab6c59582
 Port: 80/TCP
 Host Port: 0/TCP
 State: Running
 Started: Thu, 06 Dec 2018 21:25:41 +0000
 Ready: True
 Restart Count: 0
 Environment: <none>
 Mounts:
 /var/run/secrets/kubernetes.io/serviceaccount from default-token-pdn99 (ro)
Conditions:
 Type Status
 Initialized True
 Ready True
 ContainersReady True
 PodScheduled True
Volumes:
 default-token-pdn99:
 Type: Secret (a volume populated by a Secret)
 SecretName: default-token-pdn99
 Optional: false
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute for 300s
 node.kubernetes.io/unreachable:NoExecute for 300s
Events:
 Type Reason Age From Message
 ---- -
 Normal Scheduled 20s default-scheduler Successfully assigned
 default/nginxpod to ip-10-0-2-200
 Normal Pulled 19s kubelet, ip-10-0-2-200 Container image "nginx:1.11"
```

```
already present on machine
Normal Created 19s kubelet, ip-10-0-2-200 Created container
Normal Started 19s kubelet, ip-10-0-2-200 Started container
ubuntu@ip-10-0-2-200:~/pods$
```

The `kubectl` command allows you to retrieve pod metadata using the `-o` switch. The `-o` (or `--output`) switch formats the output of the get command. The output format can be *json*, *yaml*, *wide*, *name*, *template*, *template-file*, *jsonpath*, or *jsonpath-file*. The goolang template specification is also used by Docker (more info here: <http://golang.org/pkg/text/template/>) For example, to retrieve pod data in YAML, try:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl get pod nginxpod -o yaml | head

apiVersion: v1
kind: Pod
metadata:
 creationTimestamp: 2018-12-06T21:25:40Z
 name: nginxpod
 namespace: default
 resourceVersion: "28893"
 selfLink: /api/v1/namespaces/default/pods/nginxpod
 uid: 7b33b2b6-f99d-11e8-8781-02d9a858fbbc
spec:
ubuntu@ip-10-0-2-200:~/pods$
```

You can use a template to extract just the data you want.

For example, to extract just the status section's podIP value try:

```
ubuntu@ip-10-0-2-200:~/pods$ POD=$(kubectl get pod nginxpod --template=
{{.status.podIP}}) && echo $POD

10.32.0.4
ubuntu@ip-10-0-2-200:~/pods$
```

Now that we have the pod IP we can try curling our nginx server:

```
ubuntu@ip-10-0-2-200:~/pods$ curl -I $POD

HTTP/1.1 200 OK
Server: nginx/1.11.13
Date: Thu, 29 Mar 2018 05:23:37 GMT
Content-Type: text/html
Content-Length: 612
Last-Modified: Tue, 04 Apr 2017 15:01:57 GMT
Connection: keep-alive
ETag: "58e3b565-264"
Accept-Ranges: bytes
ubuntu@ip-10-0-2-200:~/pods$
```

Now that we have completed our work with the pod we can delete it:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl delete pod nginxpod

pod "nginxpod" deleted
ubuntu@ip-10-0-2-200:~/pods$
```

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl get deployment,replicaset,pods

No resources found.
ubuntu@ip-10-0-2-200:~/pods$
```

### 3. A Complex Pod

Next let's try creating a pod with a more complex specification.

Create a pod config that describes a pod with a:

- container based on an *ubuntu:14.04* image,
- with an environment variable called "MESSAGE" and a
- command that will **echo** that message to stdout
- make sure that the container is never restarted

See if you can design this specification on your own.

The pod and container spec documentation can be found here:

- Pod Spec Reference - <https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.13/#pod-v1-core>
- Container Spec Reference - <https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.13/#container-v1-core>

Create your pod when you have the configuration complete:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl create -f cpod.yaml

pod/hello created
ubuntu@ip-10-0-2-200:~/pods$
```

List your pods:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl get pod

NAME READY STATUS RESTARTS AGE
hello 0/1 Completed 0 28s
ubuntu@ip-10-0-2-200:~/pods$
```

We can verify that the container did what it was supposed to do by checking the log output of the pod using the **kubectl logs** subcommand:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl logs hello
hello world
ubuntu@ip-10-0-2-200:~/pods$
```

- Remove the hello pod

## 4. Pods and Linux Namespaces

Using the pod spec reference as a guide again, modify your nginx config (nginxpod.yaml) from step 2 so that it runs all of the containers in the pod in the host network namespace (hostNetwork). Create a new pod from your updated config.

First copy the old pod spec:

```
ubuntu@ip-10-0-2-200:~/pods$ cp nginxpod.yaml nginxpod.v2.yaml
ubuntu@ip-10-0-2-200:~/pods$
```

Next modify the copy to use the hostNetwork (use the spec reference if you need help thinking through the edits you will need to make):

```
ubuntu@ip-10-0-2-200:~/pods$ vim nginxpod.v2.yaml
...
ubuntu@ip-10-0-2-200:~/pods$
```

Run the new pod:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl create -f nginxpod.v2.yaml
pod/nginxpod created
ubuntu@ip-10-0-2-200:~/pods$
```

Now display your pod status in yaml output:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl get pod nginxpod -o yaml
apiVersion: v1
kind: Pod
metadata:
 creationTimestamp: 2018-03-29T05:44:01Z
 name: nginxpod
 namespace: default
 resourceVersion: "30554"
 selfLink: /api/v1/namespaces/default/pods/nginxpod
```

```

uid: 2eb3cf00-3314-11e8-9062-02ec217beae8
spec:
 containers:
 - image: nginx:1.11
 imagePullPolicy: IfNotPresent
 name: nginx
 ports:
 - containerPort: 80
 hostPort: 80
 protocol: TCP
 resources: {}
 terminationMessagePath: /dev/termination-log
 terminationMessagePolicy: File
 volumeMounts:
 - mountPath: /var/run/secrets/kubernetes.io/serviceaccount
 name: default-token-4kcpw
 readOnly: true
 dnsPolicy: ClusterFirst
 hostNetwork: true
 nodeName: ip-10-0-2-200
 restartPolicy: Always
 schedulerName: default-scheduler
 securityContext: {}
 serviceAccount: default
 serviceAccountName: default
 terminationGracePeriodSeconds: 30
 tolerations:
 - effect: NoExecute
 key: node.kubernetes.io/not-ready
 operator: Exists
 tolerationSeconds: 300
 - effect: NoExecute
 key: node.kubernetes.io/unreachable
 operator: Exists
 tolerationSeconds: 300
 volumes:
 - name: default-token-4kcpw
 secret:
 defaultMode: 420
 secretName: default-token-4kcpw
status:
 conditions:
 - lastProbeTime: null
 lastTransitionTime: 2018-03-29T05:44:01Z
 status: "True"
 type: Initialized
 - lastProbeTime: null
 lastTransitionTime: 2018-03-29T05:44:02Z
 status: "True"
 type: Ready
 - lastProbeTime: null
 lastTransitionTime: 2018-03-29T05:44:01Z
 status: "True"
 type: PodScheduled
 containerStatuses:
 - containerID:
 docker://fce749023c99f55bf9a58a60ebe4149fe47080eda19b438b069694e3a3a7e4e5
 image: nginx:1.11

```

```

 imageID: docker-
pullable://nginx@sha256:e6693c20186f837fc393390135d8a598a96a833917917789d63766cab6c59
582
 lastState: {}
 name: nginx
 ready: true
 restartCount: 0
 state:
 running:
 startedAt: 2018-03-29T05:44:01Z
 hostIP: 10.0.2.200
 phase: Running
 podIP: 10.0.2.200
 qosClass: BestEffort
 startTime: 2018-03-29T05:44:01Z
ubuntu@ip-10-0-2-200:~/pods$

```

If your pod is running in the host network namespace you should see that the pod (*podIP*) and host IP (*hostIP*) address are identical.

```

ubuntu@ip-10-0-2-200:~/pods$ kubectl get pod nginxpod -o yaml | grep -E "
(host|pod)IP"

 hostIP: 10.0.2.200
 podIP: 10.0.2.200
ubuntu@ip-10-0-2-200:~/pods$

```

Namespace select-ability allows you have the benefits of container deployment while still empowering infrastructure tools to see and manipulate host based networking features.

Try curling your pod using the host IP address:

```

ubuntu@ip-10-0-2-200:~/pods$ curl -I 10.0.2.200

HTTP/1.1 200 OK
Server: nginx/1.11.13
Date: Thu, 29 Mar 2018 05:44:39 GMT
Content-Type: text/html
Content-Length: 612
Last-Modified: Tue, 04 Apr 2017 15:01:57 GMT
Connection: keep-alive
ETag: "58e3b565-264"
Accept-Ranges: bytes
ubuntu@ip-10-0-2-200:~/pods$

```

- Clean up the resources (pods, etc.)

## 5. Multi-container pod

In this step we'll experiment with multi-container pods. Keep in mind that by default all containers in a pod share the same network, uts, and ipc namespace.

You can use the `--validate` switch with the `kubectl create` subcommand to verify your pod config, however the

*create* command will try to create the config regardless (work to be done here).

Create a new Pod config which:

- runs two ubuntu:14.04 containers, named hello1 and hello2
- both executing the command line “tail -f /dev/null”
- and then create it with the validate switch

You can start from an existing pod config if you like:

```
ubuntu@ip-10-0-2-200:~/pods$ cp cpod.yaml two.yaml
ubuntu@ip-10-0-2-200:~/pods$
```

Then edit the new config to meet the requirements:

```
ubuntu@ip-10-0-2-200:~/pods$ vim two.yaml
...
ubuntu@ip-10-0-2-200:~/pods$
```

Finally, create the pod:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl create -f two.yaml --validate
pod/two created
ubuntu@ip-10-0-2-200:~/pods$
```

If you got it right the first time the pod will simply run. If not you will get a descriptive error.

Issue the `kubectl get pods` command.

- How many containers are running in your new pod?
- How can you tell?

Use the `kubectl describe pod` subcommand on your new pod.

- What is the ip address of the first container?
- What is the ip address of the second container?

Shell into the first container to explore its context.

e.g. `kubectl exec -c hello1 -it two /bin/bash`

Run the `ps -ef` command inside the container.

- What processes are running in the container?
- What is the container's host name?

e.g.

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl exec -c hello1 -it two /bin/bash

root@two:/# ps -ef
UID PID PPID C STIME TTY TIME CMD
root 1 0 0 22:21 ? 00:00:00 tail -f /dev/null
root 22 0 0 22:23 pts/0 00:00:00 /bin/bash
root 36 22 0 22:23 pts/0 00:00:00 ps -ef

root@two:/# exit

exit
ubuntu@ip-10-0-2-200:~/pods$
```

Run the `ip a` command inside the container.

- What is the MAC address of eth0?
- What is the IP address

Create a file in the root directory and exit the container:

```
root@two:/# echo "Hello" > TEST

root@two:/# exit
```

Kubernetes executed our last command in the first container in the pod. We now want to open a shell into the second container. To do this we can use the `-c` switch. Exec a shell into the second container using the `-c` switch and the name of the second container:

e.g. `kubectl exec -c hello2 -it two /bin/bash`

- Is the TEST file you created previously there?
- What is the host name in this container?
- What is the MAC address of eth0?
- What is the IP address?
- Which of the following namespaces are shared across the containers in the pod?
  - User
  - Process
  - UTS (hostname)
  - Network
  - IPC
  - Mount (filesystem)

Clean up the resources (pods, etc.)

## 6. Resource Requirements



Next let's explore resource requirements. Kubernetes configs allow you to specify requested levels of memory and cpu. You can also assign limits. Requests are used when scheduling the pod to ensure that it is place on a host with enough resources free. Limits are configured in the kernel cgroups to constrain the runtime use of resources by the pod.

Create a new pod config (*limit.yaml*) like the following:

```
apiVersion: v1
kind: Pod
metadata:
 name: frontend
spec:
 containers:
 - name: db
 image: mysql
 resources:
 requests:
 memory: "64Mi"
 cpu: ".25"
 limits:
 memory: "128Mi"
 cpu: ".5"
 - name: wp
 image: wordpress
 resources:
 requests:
 memory: "64Mi"
 cpu: ".25"
 limits:
 memory: "128Mi"
 cpu: ".5"
```

This config will run a pod with the Wordpress image and the MySQL image with explicit requested resource levels and explicit resource constraints.

Before you create the pod verify the resources on your node:

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl describe $(kubectl get node -o name) | grep
-A 6 Allocated

Allocated resources:
 (Total limits may be over 100 percent, i.e., overcommitted.)
 Resource Requests Limits

 cpu 770m (38%) 0 (0%)
 memory 140Mi (3%) 340Mi (9%)
 ephemeral-storage 0 (0%) 0 (0%)
ubuntu@ip-10-0-2-200:~/pods$
```

Examine the Capacity field for your node.

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl describe $(kubectl get node -o name) -o json
| jq .status.capacity
```

```
{
 "cpu": "2",
 "ephemeral-storage": "20263528Ki",
 "hugepages-2Mi": "0",
 "memory": "4045044Ki",
 "pods": "110"
}
ubuntu@ip-10-0-2-200:~/pods$
```

If you are unfamiliar with `jq`, here is an example to list keys nested in the metadata object.

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl get $(kubectl get nodes -o name | head -1) -
o json | jq '.metadata | keys'

[
 "annotations",
 "creationTimestamp",
 "labels",
 "name",
 "resourceVersion",
 "selfLink",
 "uid"
]
ubuntu@ip-10-0-2-200:~/pods$
```

Next examine the Allocated resources section.

- Will the node be able accept the scheduled pod?

Now create the new pod and verify its construction:

```
ubuntu@ip-10-0-2-200:~$ kubectl create -f limit.yaml

pod/frontend created

ubuntu@ip-10-0-2-200:~/pods$ kubectl get pods

NAME READY STATUS RESTARTS AGE
frontend 0/2 ContainerCreating 0 11s
ubuntu@ip-10-0-2-200:~/pods$
```

Use the `kubectl describe pod frontend` or `kubectl get events | grep -i frontend` command to display the events for your new pod. You may see the image pulling. This means the Docker daemon is pulling the image in the background. You can monitor the pull by issuing the `docker pull` subcommand for the same image on the pulling host:

```
ubuntu@ip-10-0-2-200:~/pods$ docker image pull wordpress

Using default tag: latest
```

```
latest: Pulling from library/wordpress
Digest: sha256:8fd3cad0d1a9291db828ea74a7aee4cc01ff94b5ee17df493279e4d673cab56f
Status: Image is up to date for wordpress:latest
ubuntu@ip-10-0-2-200:~/pods$
```

Once the image has pulled you may see problems (Error or CrashLoopBackOff).

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl get pods

NAME READY STATUS RESTARTS AGE
frontend 1/2 Error 2 1m
ubuntu@ip-10-0-2-200:~/pods$
```

Try to diagnose and repair any issues, hint use `kubectl logs frontend -c <container name>`.

Rerun the `kubectl describe $(kubectl get node -o name)` command to redisplay your node resource usage.

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl describe $(kubectl get node -o name)

...
ubuntu@ip-10-0-2-200:~/pods$ kubectl describe $(kubectl get node -o name) | grep
-A 6 Allocated

Allocated resources:
 (Total limits may be over 100 percent, i.e., overcommitted.)
Resource Requests Limits

cpu 1270m (63%) 1 (50%)
memory 268Mi (7%) 596Mi (15%)
ephemeral-storage 0 (0%) 0 (0%)
ubuntu@ip-10-0-2-200:~/pods$
```

Can you see the impact of the new pod? Delete the pod after examining its resource usage.

```
ubuntu@ip-10-0-2-200:~/pods$ kubectl delete pod frontend

pod "frontend" deleted
ubuntu@ip-10-0-2-200:~/pods$
```

Congratulations, you have completed the lab!

## Selected Answers

Pod cpod.yaml:

```
apiVersion: v1
kind: Pod
```

```
metadata:
 name: hello
spec: # specification of the pod's contents
 restartPolicy: Never
 containers:
 - name: hellocont
 image: "ubuntu:14.04"
 env:
 - name: MESSAGE
 value: "hello world"
 command: ["/bin/sh", "-c"]
 args: ["/bin/echo \"${MESSAGE}\""]
```

Pod two.yaml:

```
apiVersion: v1
kind: Pod
metadata:
 name: two
spec:
 restartPolicy: Never
 containers:
 - name: hello1
 image: "ubuntu:14.04"
 command: ["/usr/bin/tail"]
 args: ["-f", "/dev/null"]
 - name: hello2
 image: "ubuntu:14.04"
 command: ["/usr/bin/tail"]
 args: ["-f", "/dev/null"]
```

*Copyright (c) 2013-2019 RX-M LLC, Cloud Native Consulting, all rights reserved. Licensed for use only by students in Safewrd Ventures class 18-04-19 (Max 10 students)*