

Kubernetes

Lab 4 – Deployments and Replica Sets

In this lab we will explore the nature of Kubernetes deployments and replica sets and how to work with them.

Deployments

A deployment provides declarative updates for pods and replica sets. You describe the desired state in a deployment object, and the deployment controller will change the actual state to the desired state at a controlled rate for you. You can define deployments to create new resources, or replace existing ones by new ones. Typical uses:

- bring up a replica set and (indirectly) its pods
- capturing the results and status of a deployment
- updating an existing deployment to recreate pods with a new image (rolling updates)
- rolling back to an earlier deployment revision if the current deployment isn't stable
- pausing and resuming a deployment

ReplicaSets

Replica sets (RS) supersede the older replication controller (RC) resource type. Replica sets support the set-based selectors as well as equality-based selector requirements (RCs only supported equality.) While replica sets can be used independently, they are mainly used by deployments as a mechanism to orchestrate pod creation, deletion, and updates. When you use deployments you don't have to worry about managing the replica sets that they create; deployments own and manage their replica sets.

ReplicaSets ensure that a specified number of pod "replicas" are running at all times. If there are too many, it will kill some. If there are too few, it will start more. Unlike in the case where a user directly created pods, a ReplicaSet replaces pods that are deleted or terminated for any reason, such as in the case of node failure or disruptive node maintenance (e.g. a kernel upgrade, etc.)

For this reason the Kubernetes team recommends that you use a Deployment/ReplicaSet even if your application requires only a single pod. ReplicaSets are like process supervisors in many ways but monitor processes on multiple nodes at once. A ReplicaSet delegates local container restarts to some agent on the node (e.g., Kubelet or Docker.)

A ReplicaSet is only appropriate for pods with *RestartPolicy = Always* (if the RestartPolicy is not set, the default value is *Always*.) A ReplicaSet will refuse to instantiate any pod that has a different restart policy.

A ReplicaSet will never terminate on its own, but it isn't expected to be as long-lived as services. Services may be composed of pods controlled by multiple ReplicaSets, and it is expected that many ReplicaSets may be created and destroyed over the lifetime of a service (for instance, to perform an update of pods that run the service.) Both services themselves and their clients should remain oblivious to the ReplicaSets that maintain the pods of the services.

Now to create some Deployments/ReplicaSets.

1. A Simple Deployment

As a first exploratory step lets create a simple deployment which stands up three nginx pods. Create a config file similar to the following to accomplish this task:

```
ubuntu@ip-10-0-2-200:~$ cd ~
ubuntu@ip-10-0-2-200:~$ mkdir dep
ubuntu@ip-10-0-2-200:~$ cd dep
ubuntu@ip-10-0-2-200:~/dep$
ubuntu@ip-10-0-2-200:~/dep$ vi mydep.yaml
ubuntu@ip-10-0-2-200:~/dep$ cat mydep.yaml

apiVersion: apps/v1
kind: Deployment
metadata:
  name: website
  labels:
    bu: sales
spec:
  replicas: 3
  selector:
    matchLabels:
      appname: webserver
      targetenv: demo
  template:
    metadata:
      labels:
        appname: webserver
        targetenv: demo
    spec:
      containers:
      - name: podweb
        image: nginx:1.7.9
        ports:
        - containerPort: 80
ubuntu@ip-10-0-2-200:~/dep$
```

Deployments were promoted to the apps/v1 API in K8s 1.9 but were added to Kubernetes 1.2 and are the go forward solution for deploying replicated pods. The spec for Replication Controllers (part of the v1 API) is almost the same as the spec for Deployments though deployments add a few key features such as the ability to specify upgrades declaratively. The specification for Deployments can be found [here](https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.13/#deployment-v1-apps):

<https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.13/#deployment-v1-apps>

Now create the Deployment using the `kubectl create` subcommand and verify that the Deployment, its ReplicaSet and pods are up with the `get` subcommand:

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl create -f mydep.yaml

deployment.apps/website created
```

```
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get deploy,rs,pods
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.extensions/website	3	3	3	3	13s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.extensions/website-6dc99878b	3	3	3	13s

NAME	READY	STATUS	RESTARTS	AGE
pod/website-6dc99878b-4w9tv	1/1	Running	0	13s
pod/website-6dc99878b-8wh9l	1/1	Running	0	13s
pod/website-6dc99878b-q5jqx	1/1	Running	0	13s

```
ubuntu@ip-10-0-2-200:~/dep$
```

While everything appears to be running we can verify that there are no scheduling cycles or fail/restart activities by examining the system events. We have viewed resource specific events in the past using the `kubectl describe` subcommand. This time we'll use the `kubectl get events` subcommand to view cluster wide events:

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get events --sort-by='{.lastTimestamp}' |  
grep website
```

55s	Normal	ScalingReplicaSet	Deployment	Scaled up replica set website-6dc99878b to 3
55s	Normal	SuccessfulCreate	ReplicaSet	Created pod: website-6dc99878b-q5jqx
54s	Normal	SuccessfulCreate	ReplicaSet	Created pod: website-6dc99878b-8wh9l
54s	Normal	SuccessfulCreate	ReplicaSet	Created pod: website-6dc99878b-4w9tv
54s	Normal	Scheduled	Pod	Successfully assigned default/website-6dc99878b-8wh9l to ip-10-0-2-200
54s	Normal	Scheduled	Pod	Successfully assigned default/website-6dc99878b-q5jqx to ip-10-0-2-200
54s	Normal	Scheduled	Pod	Successfully assigned default/website-6dc99878b-4w9tv to ip-10-0-2-200

```
ubuntu@ip-10-0-2-200:~/dep$
```

Checking the event log occasionally will help you identify normal cluster patterns and make it possible for you to spot anomalies more easily when debugging.

The replica set began with a scale of 3, causing 3 instances of the pod template to get scheduled. Replica sets ensure that some number of instances of the pod template are always running. Try deleting a pod (use the pod name displayed by `kubectl get pods`).

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl delete pod website-6dc99878b-q5jqx  
  
pod "website-6dc99878b-q5jqx" deleted
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
website-6dc99878b-4w9tv	1/1	Running	0	33s
website-6dc99878b-8wh9l	1/1	Running	0	33s
website-6dc99878b-l9rmz	1/1	Running	0	22s

You might ask: "why would kubernetes let someone delete the pod if it will just restart it?". There are many reasons you might want to delete a given pod. Perhaps it has problems and you want to generate a new replacement. Perhaps the current node has problems and you want Kubernetes to reschedule this particular pod somewhere else. To actually terminate our pod permanently we must delete the deployment, the deployment controls the replica set, the replica set controls the pods.

When many resources are running on a cluster it can be advantageous to restrict output to a certain set of resources. The Kubernetes labeling system makes this easy. The `-l` switch can be used with the `kubectl get` subcommand to filter output by label.

Try listing all pods:

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
website-6dc99878b-4w9tv	1/1	Running	0	85s
website-6dc99878b-8wh9l	1/1	Running	0	85s
website-6dc99878b-l9rmz	1/1	Running	0	74s

```
ubuntu@ip-10-0-2-200:~/dep$
```

Now try filtering by the "appname" label key we assigned to all of our pods in the pod template metadata:

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get pods -l appname
```

NAME	READY	STATUS	RESTARTS	AGE
website-6dc99878b-4w9tv	1/1	Running	0	97s
website-6dc99878b-8wh9l	1/1	Running	0	97s
website-6dc99878b-l9rmz	1/1	Running	0	86s

```
ubuntu@ip-10-0-2-200:~/dep$
```

You can also filter by key and value:

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get pods -l appname=webserver
```

NAME	READY	STATUS	RESTARTS	AGE
website-6dc99878b-4w9tv	1/1	Running	0	113s
website-6dc99878b-8wh9l	1/1	Running	0	113s
website-6dc99878b-l9rmz	1/1	Running	0	102s

```
ubuntu@ip-10-0-2-200:~/dep$
```

You can filter by pod name:

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get $(kubectl get pods -o name | head -1)

NAME                                READY   STATUS    RESTARTS   AGE
website-6dc99878b-4w9tv            1/1     Running   0           2m6s
ubuntu@ip-10-0-2-200:~/dep$
```

Our pod has labels we have added and the Kubernetes infrastructure may add labels as well:

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl describe $(kubectl get pods -o name | head -1) | grep -A2 -i label

Labels:      appname=webserver
             pod-template-hash=6dc99878b
             targetenv=demo
ubuntu@ip-10-0-2-200:~/dep$
```

Unfortunately describe doesn't allow for JSON output. Good news, though, `get` does.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get $(kubectl get pods -o name | head -1) -o json | jq .metadata.labels
```

```
{
  "appname": "webserver",
  "pod-template-hash": "6dc99878b",
  "targetenv": "demo"
}
```

```
ubuntu@ip-10-0-2-200:~/dep$
```

- Why do each of the filters above work or not work?
- Enter a command to display all of the pods with either the “demo” or “prod” value for targetenv
- Find all pods other than those with the “demo” or “prod” value for targetenv
- Enter a command to display all of the pods with either the “demo” or “prod” value for targetenv and the appname key set to webserver

2. Checking status of a Deployment

We have seen previously how to check the status of a deployment.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get deploy

NAME      DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
website   3          3          3             3           2m
ubuntu@ip-10-0-2-200:~/dep$
```

Now we take an slightly more application-centric view.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl rollout status deploy/website  
  
deployment "website" successfully rolled out  
ubuntu@ip-10-0-2-200:~/dep$
```

Rollouts are used to update a given set of Pods, the ones controlled by this Deployment's replica set. It reports success when all the currently deployed Pods match what is expected in the current deployment. In k8s technical terms these conditions are all true:

- `.status.observedGeneration >= .metadata.generation`
- `.status.updatedReplicas == .spec.replicas`
- `.spec.availableReplicas >= minimum required`

3. Updating a Deployment

We are using nginx 1.7.9 in our example, lets update to 1.9.1.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl set image deploy/website podweb=nginx:1.9.1 --record  
  
deployment.extensions/website image updated  
ubuntu@ip-10-0-2-200:~/dep$
```

Alternative is to use `kubectl edit deployment/website`

Check the status of the rollout (if you're not fast you may not see these updates):

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl rollout status deploy/website  
  
Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have  
been updated...  
Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have  
been updated...  
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have  
been updated...  
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have  
been updated...  
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have  
been updated...  
Waiting for deployment "website" rollout to finish: 2 out of 3 new replicas have  
been updated...  
Waiting for deployment "website" rollout to finish: 1 old replicas are pending  
termination...  
Waiting for deployment "website" rollout to finish: 1 old replicas are pending  
termination...  
deployment "website" successfully rolled out  
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get deploy/website
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
website	3	3	3	3	5m

```
ubuntu@ip-10-0-2-200:~/dep$
```

Look at the Replica Sets & Pods

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get rs,pod
```

NAME	DESIRED	CURRENT	READY	AGE
replicaset.extensions/website-654d96bc8d	3	3	3	51s
replicaset.extensions/website-6dc99878b	0	0	0	5m13s

NAME	READY	STATUS	RESTARTS	AGE
pod/website-654d96bc8d-2r225	1/1	Running	0	43s
pod/website-654d96bc8d-ckqlg	1/1	Running	0	45s
pod/website-654d96bc8d-l8jr8	1/1	Running	0	51s

```
ubuntu@ip-10-0-2-200:~/dep$
```

By describing the deployment we can inspect the events that occurred during the rollout:

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl describe deploy/website | grep -A 10 Events
```

```
Events:
  Type    Reason             Age   From              Message
  ----    -
  Normal  ScalingReplicaSet  5m32s deployment-controller Scaled up replica set website-6dc99878b to 3
  Normal  ScalingReplicaSet  70s   deployment-controller Scaled up replica set website-654d96bc8d to 1
  Normal  ScalingReplicaSet  64s   deployment-controller Scaled down replica set website-6dc99878b to 2
  Normal  ScalingReplicaSet  64s   deployment-controller Scaled up replica set website-654d96bc8d to 2
  Normal  ScalingReplicaSet  62s   deployment-controller Scaled down replica set website-6dc99878b to 1
  Normal  ScalingReplicaSet  62s   deployment-controller Scaled up replica set website-654d96bc8d to 3
  Normal  ScalingReplicaSet  60s   deployment-controller Scaled down replica set website-6dc99878b to 0
ubuntu@ip-10-0-2-200:~/dep$
```

Note that the rollout was a smooth transition from one set of Pods controlled by our original ReplicaSet `website-6dc99878b` to our second set of Pods controlled by the RS `website-6dc99878b`.

4. Manually rolling back a deployment

Lets manually revert back to nginx 1.7.9 and check the status.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl set image deploy/website podweb=nginx:1.7.9 --record
```

```
deployment.extensions/website image updated
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl rollout status deploy/website
...
Waiting for rollout to finish: 1 old replicas are pending termination...
Waiting for rollout to finish: 1 old replicas are pending termination...
deployment "website" successfully rolled out
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
website-654d96bc8d	0	0	0	2m49s
website-6dc99878b	3	3	3	7m11s

```
ubuntu@ip-10-0-2-200:~/dep$
```

Notice which deployment (NAME) is being used.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
website-6dc99878b-98wqm	1/1	Running	0	59s
website-6dc99878b-cfd5j	1/1	Running	0	55s
website-6dc99878b-mr684	1/1	Running	0	57s

```
ubuntu@ip-10-0-2-200:~/dep$
```

Confirm your observations once again in the event log.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl describe deploy/website | grep -A 15 Events
```

```
Events:
  Type    Reason             Age          From              Message
  ----    -
  Normal  ScalingReplicaSet  3m26s       deployment-controller Scaled
up replica set website-654d96bc8d to 1
  Normal  ScalingReplicaSet  3m20s       deployment-controller Scaled
down replica set website-6dc99878b to 2
  Normal  ScalingReplicaSet  3m20s       deployment-controller Scaled
up replica set website-654d96bc8d to 2
  Normal  ScalingReplicaSet  3m18s       deployment-controller Scaled
down replica set website-6dc99878b to 1
```



```

Normal ScalingReplicaSet 3m18s      deployment-controller Scaled
up replica set website-654d96bc8d to 3
Normal ScalingReplicaSet 3m16s      deployment-controller Scaled
down replica set website-6dc99878b to 0
Normal ScalingReplicaSet 72s        deployment-controller Scaled
up replica set website-6dc99878b to 1
Normal ScalingReplicaSet 70s        deployment-controller Scaled
down replica set website-654d96bc8d to 2
Normal ScalingReplicaSet 68s (x2 over 7m48s) deployment-controller Scaled
up replica set website-6dc99878b to 3
Normal ScalingReplicaSet 66s (x3 over 70s) deployment-controller
(combined from similar events): Scaled down replica set website-654d96bc8d to 0
ubuntu@ip-10-0-2-200:~/dep$

```

5. Checking rollout history of a Deployment

We can use the *rollout history* subcommand to see what we have been doing to trigger these rollouts

```

ubuntu@ip-10-0-2-200:~/dep$ kubectl rollout history deploy/website

deployment.extensions/website
REVISION  CHANGE-CAUSE
2         kubectl set image deploy/website podweb=nginx:1.9.1 --record=true
3         kubectl set image deploy/website podweb=nginx:1.7.9 --record=true
ubuntu@ip-10-0-2-200:~/dep$

```

Take a detailed look at a previous deployment version.

```

ubuntu@ip-10-0-2-200:~/dep$ kubectl rollout history deploy/website --revision=2

deployments "website" with revision #2
Pod Template:
  Labels:      appname=webserver
              pod-template-hash=654d96bc8d
              targetenv=demo
  Annotations: kubernetes.io/change-cause: kubectl set image deploy/website
podweb=nginx:1.9.1 --record=true
  Containers:
   podweb:
    Image:      nginx:1.9.1
    Port:       80/TCP
    Host Port:  0/TCP
    Environment: <none>
    Mounts:      <none>
    Volumes:      <none>

ubuntu@ip-10-0-2-200:~/dep$

```

6. Rolling back to a previous Deployment

Confirm the current version of a container is 1.7.9.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get pods -o json | jq
.items[0].spec.containers[0].image -r

nginx:1.7.9
ubuntu@ip-10-0-2-200:~/dep$
```

Revert to previous version/revision.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl rollout undo deploy/website

deployment.extensions/website
ubuntu@ip-10-0-2-200:~/dep$
```

Alternative to above is `kubectl rollout undo deployment/website --to-revision=2`

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get deploy/website
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
website	3	3	3	3	8m

```
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl describe deploy/website | grep -A 15 Events
```

```
Events:
  Type      Reason      Age      From
  Message
  ----      -
  --
    Normal ScalingReplicaSet  4m56s      deployment-controller
Scaled down replica set website-6dc99878b to 2
    Normal ScalingReplicaSet  4m54s      deployment-controller
Scaled down replica set website-6dc99878b to 1
    Normal ScalingReplicaSet  4m54s      deployment-controller
Scaled up replica set website-654d96bc8d to 3
    Normal ScalingReplicaSet  2m48s      deployment-controller
Scaled up replica set website-6dc99878b to 1
    Normal ScalingReplicaSet  2m46s      deployment-controller
Scaled down replica set website-654d96bc8d to 2
    Normal ScalingReplicaSet  2m44s (x2 over 9m24s)  deployment-controller
Scaled up replica set website-6dc99878b to 3
    Normal ScalingReplicaSet  30s (x2 over 5m2s)    deployment-controller
Scaled up replica set website-654d96bc8d to 1
    Normal DeploymentRollback  30s      deployment-controller
Rolled back deployment "website" to revision 2
    Normal ScalingReplicaSet  29s (x2 over 4m56s)    deployment-controller
Scaled up replica set website-654d96bc8d to 2
    Normal ScalingReplicaSet  27s (x6 over 2m46s)    deployment-controller
(combined from similar events): Scaled up replica set website-654d96bc8d to 3
```

```
Normal ScalingReplicaSet 25s (x2 over 4m52s) deployment-controller
Scaled down replica set website-6dc99878b to 0
ubuntu@ip-10-0-2-200:~/dep$
```

Note the unique event in the log for the rollback: `DeploymentRollback`

Confirm the container image version has been reverted to 1.9.1:

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get pods -o json | jq
.items[0].spec.containers[0].image -r

nginx:1.9.1
ubuntu@ip-10-0-2-200:~/dep$
```

7. Pausing and resuming a Deployment

In a larger installation, we may be deploying dozens of pods. For our small test it is hard to pause in time, so we chain the commands to hopefully catch it in the act.

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl set image deploy/website podweb=nginx:1.7.9;
kubectl rollout pause deploy/website

deployment.extensions/website image updated
deployment.extensions/website paused
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get rs

NAME                DESIRED   CURRENT   READY   AGE
website-654d96bc8d   3         3         3       6m22s
website-6dc99878b    1         1         1       10m
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl rollout status deploy/website

Waiting for deployment "website" rollout to finish: 1 out of 3 new replicas have
been updated...
^C
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl rollout resume deploy/website

deployment.extensions/website resumed
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl rollout status deploy/website

deployment "website" successfully rolled out
ubuntu@ip-10-0-2-200:~/dep$
```

```
ubuntu@ip-10-0-2-200:~/dep$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
website-654d96bc8d	0	0	0	7m28s
website-6dc99878b	3	3	3	11m

```
ubuntu@ip-10-0-2-200:~/dep$
```

Delete your deployment.

8. Health Checks

In this step we will create a pod with a health check. Enter and run the following config (*hc.yaml*):

```
ubuntu@ip-10-0-2-200:~/dep$ cd ~
ubuntu@ip-10-0-2-200:~$ mkdir hc
ubuntu@ip-10-0-2-200:~$ cd hc
ubuntu@ip-10-0-2-200:~/hc$ vi hc.yaml
ubuntu@ip-10-0-2-200:~/hc$ cat hc.yaml

apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      name: nginx
  template:
    metadata:
      labels:
        name: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:latest
        ports:
        - containerPort: 80
        livenessProbe: # An HTTP health check
          httpGet:
            path: /
```

```
port: 80
initialDelaySeconds: 30
timeoutSeconds: 1
ubuntu@ip-10-0-2-200:~/hc$
```

Now run the deployment:

```
ubuntu@ip-10-0-2-200:~/hc$ kubectl create -f hc.yaml

deployment.apps/nginx created
ubuntu@ip-10-0-2-200:~/hc$
```

View your deployment:

```
ubuntu@ip-10-0-2-200:~/hc$ kubectl get deploy,rs,pods
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.extensions/nginx	3	3	3	3	14s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.extensions/nginx-5fb57bb9c9	3	3	3	14s

NAME	READY	STATUS	RESTARTS	AGE
pod/nginx-5fb57bb9c9-b6vh2	1/1	Running	0	14s
pod/nginx-5fb57bb9c9-mf599	1/1	Running	0	14s
pod/nginx-5fb57bb9c9-nxnz2	1/1	Running	0	14s

```
ubuntu@ip-10-0-2-200:~/hc$
```

Note that our nginx service listens on port 80 and responds normally to requests for "/", so our health check is passing.

To trigger the health check repair logic, we need to simulate an error condition. By forcing nginx to report a 404, the `HttpGet` livenessProbe will fail. We can do this by deleting the nginx configuration file in the nginx container.

Display the events for the first pod in the set:

```
ubuntu@ip-10-0-2-200:~/hc$ kubectl get events --sort-by='{.lastTimestamp}' \
| grep $(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}')
```

34s	Normal	Scheduled	Pod	Successfully assigned
default/nginx-5fb57bb9c9-b6vh2 to ip-10-0-2-200				
34s	Normal	SuccessfulCreate	ReplicaSet	Created pod: nginx-
5fb57bb9c9-b6vh2				

```
ubuntu@ip-10-0-2-200:~/hc$
```

The status is good.

Now lets tell the nginx in the first pod to stop serving the root IRI by deleting the nginx default config.

```
ubuntu@ip-10-0-2-200:~/hc$ kubectl exec -it $(kubectl get pods -o name | head -1
```

```
| awk -F '/' '{print $2}') \
-- sh -c "rm /etc/nginx/conf.d/default.conf && nginx -s reload"

2018/03/30 00:06:59 [notice] 15#15: signal process started
ubuntu@ip-10-0-2-200:~/hc$
```

Now redisplay the events for the pod:

```
ubuntu@ip-10-0-2-200:~/hc$ kubectl get events --sort-by='{.lastTimestamp}' \
| grep $(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}')
```

65s	Normal	Scheduled	Pod	Successfully assigned
default/nginx-5fb57bb9c9-b6vh2 to ip-10-0-2-200				
65s	Normal	SuccessfulCreate	ReplicaSet	Created pod: nginx-
5fb57bb9c9-45x6s				

```
ubuntu@ip-10-0-2-200:~/hc$
```

What happened?

Events reported by the event stream are not as granular as those provided by the describe, try it:

```
ubuntu@ip-10-0-2-200:~/hc$ $ kubectl describe pod \
$(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}') | grep -A 15
Events
```

Type	Reason	Age	From	Message
Normal	Scheduled	4m23s	default-scheduler	Successfully assigned default/nginx-5fb57bb9c9-b6vh2 to ip-10-0-2-200
Warning	Unhealthy	2m56s (x3 over 3m16s)	kubelet, ip-10-0-2-200	Liveness probe failed: Get http://10.32.0.5:80/: dial tcp 10.32.0.5:80: connect: connection refused
Normal	Pulling	2m55s (x2 over 4m22s)	kubelet, ip-10-0-2-200	pulling image "nginx:latest"
Normal	Pulled	2m55s (x2 over 4m21s)	kubelet, ip-10-0-2-200	Successfully pulled image "nginx:latest"
Normal	Created	2m55s (x2 over 4m20s)	kubelet, ip-10-0-2-200	Created container
Normal	Started	2m55s (x2 over 4m20s)	kubelet, ip-10-0-2-200	Started container
Normal	Killing	2m55s	kubelet, ip-10-0-2-200	Killing container with id docker://nginx:Container failed liveness probe.. Container will be killed and recreated.

```
ubuntu@ip-10-0-2-200:~/hc$
```

As you can see the Liveness probe is now failing. The nginx container in the pod was created, started, found unhealthy, killed, created and started again.

Remove the related resources.

```
ubuntu@ip-10-0-2-200:~/jobs$ kubectl delete deploy/nginx  
deployment.extensions "nginx" deleted  
ubuntu@ip-10-0-2-200:~/jobs$
```

9. Creating a Job

In a previous lab we saw that running a pod standalone works but without an RS the pod will not restart if it crashes. Unfortunately, if we run a batch job in a pod with an RS and the pod completes the task, the RS will start the pod again.

What if we want a pod that runs only once, however, if it or the node it is running on fails before the pod completes successfully, we want the pod to be started again until it does complete successfully. Kubernetes provides a Job type for this scenario.

A Job is like an RC/RS that ensures that a pod runs once to completion. Imagine we want to calculate Pi. Not twice, not half of a time, but precisely once. A job would be the perfect way to run a container that calculates Pi. Enter this sample job config to compute Pi:

```
ubuntu@ip-10-0-2-200:~/hc$ cd ~  
ubuntu@ip-10-0-2-200:~$ mkdir jobs  
ubuntu@ip-10-0-2-200:~$ cd jobs/  
ubuntu@ip-10-0-2-200:~/jobs$ vim myjob.yaml  
ubuntu@ip-10-0-2-200:~/jobs$ cat myjob.yaml  
  
apiVersion: batch/v1  
kind: Job  
metadata:  
  name: pi  
spec:  
  template:  
    metadata:  
      name: pi  
    spec:  
      containers:  
      - name: pi  
        image: perl  
        command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]  
        restartPolicy: Never  
ubuntu@ip-10-0-2-200:~/jobs$
```

The config uses apiVersion "batch/v1". The kind of object we will create is a Job. The Job will have the name "pi", as per the metadata. The template for the pod the Job we'll create must have a name pi.

The spec for the pod uses a single perl container which will run the command that computes pi. We also set the restart policy to Never.

Now try running your Job:

```
ubuntu@ip-10-0-2-200:~/jobs$ kubectl create -f myjob.yaml

job.batch/pi created
ubuntu@ip-10-0-2-200:~/jobs$
```

Examine the job:

```
ubuntu@ip-10-0-2-200:~/jobs$ kubectl get deploy,rs,pods,job

NAME                READY   STATUS             RESTARTS   AGE
pod/pi-xzx2h        0/1     ContainerCreating   0           20s

NAME                COMPLETIONS   DURATION   AGE
job.batch/pi        0/1           20s        20s
ubuntu@ip-10-0-2-200:~/jobs$
```

```
ubuntu@ip-10-0-2-200:~/jobs$ kubectl get deploy,rs,pods,job

NAME                READY   STATUS      RESTARTS   AGE
pod/pi-xzx2h        0/1     Completed   0           30s

NAME                COMPLETIONS   DURATION   AGE
job.batch/pi        1/1           27s        30s
ubuntu@ip-10-0-2-200:~/jobs$
```

```
ubuntu@ip-10-0-2-200:~/jobs$ kubectl describe job/pi

Name:                pi
Namespace:           default
Selector:             controller-uid=59c75687-f9cb-11e8-8781-02d9a858fbbc
Labels:              controller-uid=59c75687-f9cb-11e8-8781-02d9a858fbbc
                    job-name=pi
Annotations:         <none>
Parallelism:         1
Completions:         1
Start Time:          Fri, 07 Dec 2018 02:54:01 +0000
Completed At:        Fri, 07 Dec 2018 02:54:28 +0000
Duration:            27s
Pods Statuses:       0 Running / 1 Succeeded / 0 Failed
Pod Template:
  Labels:  controller-uid=59c75687-f9cb-11e8-8781-02d9a858fbbc
          job-name=pi
  Containers:
    pi:
      Image:  perl
      Port:   <none>
      Command:
        perl
        -Mbignum=bpi
        -wle
        print bpi(2000)
```



```

Environment: <none>
Mounts:      <none>
Volumes:     <none>
Events:
  Type        Reason              Age   From                Message
  ----        -
  Normal      SuccessfulCreate    52s   job-controller      Created pod: pi-xzx2h
ubuntu@ip-10-0-2-200:~/jobs$

```

The `kubectl create` subcommand processes the job request and runs our pod. Displaying the Job description shows us the name of the pod that ran the Job. We can now dump the logs for the pod to see the result:

```

ubuntu@ip-10-0-2-200:~/jobs$ kubectl logs $(kubectl get jobs -o name)

3.1415926535897932384626433832795028841971693993751058209749445923078164062862089
986280348253421170679821480865132823066470938446095505822317253594081284811174502
841027019385211055596446229489549303819644288109756659334461284756482337867831652
712019091456485669234603486104543266482133936072602491412737245870066063155881748
815209209628292540917153643678925903600113305305488204665213841469519415116094330
572703657595919530921861173819326117931051185480744623799627495673518857527248912
279381830119491298336733624406566430860213949463952247371907021798609437027705392
171762931767523846748184676694051320005681271452635608277857713427577896091736371
787214684409012249534301465495853710507922796892589235420199561121290219608640344
181598136297747713099605187072113499999983729780499510597317328160963185950244594
553469083026425223082533446850352619311881710100031378387528865875332083814206171
776691473035982534904287554687311595628638823537875937519577818577805321712268066
130019278766111959092164201989380952572010654858632788659361533818279682303019520
353018529689957736225994138912497217752834791315155748572424541506959508295331168
617278558890750983817546374649393192550604009277016711390098488240128583616035637
076601047101819429555961989467678374494482553797747268471040475346462080466842590
694912933136770289891521047521620569660240580381501935112533824300355876402474964
732639141992726042699227967823547816360093417216412199245863150302861829745557067
498385054945885869269956909272107975093029553211653449872027559602364806654991198
818347977535663698074265425278625518184175746728909777727938000816470600161452491
921732172147723501414419735685481613611573525521334757418494684385233239073941433
345477624168625189835694855620992192221842725502542568876717904946016534668049886
272327917860857843838279679766814541009538837863609506800642251252051173929848960
841284886269456042419652850222106611863067442786220391949450471237137869609563643
719172874677646575739624138908658326459958133904780275898
ubuntu@ip-10-0-2-200:~/jobs$

```

By default, a Job is complete when one Pod runs to successful completion. You can also specify that this needs to happen multiple times by specifying Job spec key `"completions"` with a value greater than 1. You can suggest how many pods should run concurrently by setting Job spec key `"parallelism"` to the number of pods you would like to have running concurrently (the value defaults to `"completions"`.) The parallelism key is just a hint and the Job may run fewer or more concurrent pods.

Jobs are complementary to Deployments. A Deployment manages pods which are not expected to terminate (e.g. web servers,) and a Job manages pods that are expected to terminate (e.g. batch jobs.)

When you are finished exploring remove the Job:

```
ubuntu@ip-10-0-2-200:~/jobs$ kubectl delete job pi
```

```
job.batch "pi" deleted  
ubuntu@ip-10-0-2-200:~/jobs$
```

```
ubuntu@ip-10-0-2-200:~/jobs$ kubectl get deploy,rs,pods,job
```

```
No resources found.  
ubuntu@ip-10-0-2-200:~/jobs$
```

Congratulations you have completed the lab!

Copyright (c) 2013-2019 RX-M LLC, Cloud Native Consulting, all rights reserved. Licensed for use only by students in Safewrd Ventures class 18-04-19 (Max 10 students)